

Exploratory Data Analysis

```
In [1]: # Lecture imports / dependencies
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import pandas as pd
import seaborn as sns
sns.set(style="ticks")
from sklearn.feature_extraction.text import CountVectorizer
from skimage.io import imread, imshow
```

Typical steps of ML/DM

1. Identify question / task
2. Collect data
3. Clean and preprocess data
4. Exploratory data analysis (EDA)
5. Feature and model selection
6. Train model
7. Evaluate and communicate results
8. Deploy working system

(but not necessarily in this order...)

Today we'll discuss steps (3) and (4)

What does data look like?

Often, it is tabular (but certainly not always!).

```
In [6]: titanic = sns.load_dataset("titanic")
titanic.head()
```

```
Out[6]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	embark_town	alive	alone
0	0	3	male	22.0	1	0	7.2500	S	Third	man	True	NaN	Southampton	no	False
1	1	1	female	38.0	1	0	71.2833	C	First	woman	False	C	Cherbourg	yes	False
2	1	3	female	26.0	0	0	7.9250	S	Third	woman	False	NaN	Southampton	yes	True
3	1	1	female	35.0	1	0	53.1000	S	First	woman	False	C	Southampton	yes	False
4	0	3	male	35.0	0	0	8.0500	S	Third	man	True	NaN	Southampton	no	True

```
In [16]: titanic.dtypes
```

```
Out[16]: survived      int64
pclass      int64
sex         object
age         float64
sibsp       int64
parch       int64
fare        float64
embarked    object
class       category
who         object
adult_male  bool
deck        category
embark_town object
alive       object
alone       bool
dtype: object
```

```
In [17]: titanic.groupby("deck").size()
```

```
Out[17]: deck
A      15
B      47
C      59
D      33
E      32
F      13
G       4
dtype: int64
```

- Each row is an **object** (or training example, or sample)
- Each column is a **feature** (or variable, covariate).

Types of features

- Categorical (e.g. `survived` , `embark_town`)
- Numerical (e.g. `age`, `fare`)
- Some are more ambiguous, like `pclass` : is this categorical or numerical?


Converting types:

- Many of our methods are meant to work with numerical features.
- We can convert categorical to numerical.

```
In [11]: pd.get_dummies(titanic, columns=["class"]).head()
```

```
Out[11]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	who	adult_male	deck	embark_town	alive	alone	class_First	class_Second
0	0	3	male	22.0	1	0	7.2500	S	man	True	NaN	Southampton	no	False	0	0
1	1	1	female	38.0	1	0	71.2833	C	woman	False	C	Cherbourg	yes	False	1	1
2	1	3	female	26.0	0	0	7.9250	S	woman	False	NaN	Southampton	yes	True	0	0
3	1	1	female	35.0	1	0	53.1000	S	woman	False	C	Southampton	yes	False	1	1
4	0	3	male	35.0	0	0	8.0500	S	man	True	NaN	Southampton	no	True	0	0



```
In [12]: titanic.shape
```

```
Out[12]: (891, 15)
```

```
In [111]: titanic.describe()
```

```
Out[111]:
```

	survived	pclass	age	sibsp	parch	fare
count	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

If we do this for all our features, we can now interpret objects as points in space.

```
In [112]: titanic_num = pd.get_dummies(titanic, columns=["sex","embarked","class","who","adult_male","deck","embark_town","alive"])
titanic_num.shape
```

Out[112]: (891, 33)

```
In [113]: titanic_num.head()
```

Out[113]:

	survived	pclass	age	sibsp	parch	fare	sex_female	sex_male	embarked_C	embarked_Q	...	deck_E	deck_F	deck_G	embark_town_Ch
0	0	3	22.0	1	0	7.2500	0	1	0	0	...	0	0	0	
1	1	1	38.0	1	0	71.2833	1	0	1	0	...	0	0	0	
2	1	3	26.0	0	0	7.9250	1	0	0	0	...	0	0	0	
3	1	1	35.0	1	0	53.1000	1	0	0	0	...	0	0	0	
4	0	3	35.0	0	0	8.0500	0	1	0	0	...	0	0	0	

5 rows × 33 columns



```
In [114]: titanic_num.dtypes
```

```
Out[114]: survived                int64
pclass                            int64
age                               float64
sibsp                             int64
parch                             int64
fare                              float64
sex_female                        uint8
sex_male                          uint8
embarked_C                        uint8
embarked_Q                        uint8
embarked_S                        uint8
class_First                       uint8
class_Second                      uint8
class_Third                       uint8
who_child                         uint8
who_man                           uint8
who_woman                         uint8
adult_male_False                  uint8
adult_male_True                   uint8
deck_A                           uint8
deck_B                           uint8
deck_C                           uint8
deck_D                           uint8
deck_E                           uint8
deck_F                           uint8
deck_G                           uint8
embark_town_Cherbourg             uint8
embark_town_Queenstown            uint8
embark_town_Southampton           uint8
alive_no                          uint8
alive_yes                         uint8
alone_False                       uint8
alone_True                        uint8
dtype: object
```

- So we now have 891 objects and 33 features.
- In other words, each object is a point in 33-dimensional space.
- This is why multivariable calculus is a prerequisite.

Other feature types: text data

```
In [36]: text = "The University of British Columbia (UBC) is a public research university with campuses and facilities in Briti
```



One approach: **bag of words** features.

```
# clustering of google documents  
# checking if your mail is spam or not  
# analyzing the reviews of customer on your product (positive or negative)
```

```
In [56]: cv = CountVectorizer()  
feat = cv.fit_transform([text])
```

```
In [57]: for word, idx in cv.vocabulary_.items():  
         print("%-14s%d" % (word, feat[0,idx]))
```

the	1
university	2
of	1
british	2
columbia	2
ubc	1
is	1
public	1
research	1
with	1
campuses	1
and	1
facilities	1
in	1
canada	1

- Bag of words ignores the order of words but still can work well.

Other feature types: images

```
In [ ]: # facebook use images to tag friends and people (object recognition)
```

```
In [64]: img = imread("Mufic.jpg")  
plt.xticks([])  
plt.yticks([])  
imshow(img);
```



Photo credit: [Wikipedia: UBC](https://en.wikipedia.org/wiki/University_of_British_Columbia#/media/File:Irving_K._Barber_Library.jpg) (https://en.wikipedia.org/wiki/University_of_British_Columbia#/media/File:Irving_K._Barber_Library.jpg) by [CjayD](https://www.flickr.com/people/85424459@N08/) (<https://www.flickr.com/people/85424459@N08/>), [CC BY 2.0](http://creativecommons.org/licenses/by/2.0) (<http://creativecommons.org/licenses/by/2.0>).

```
In [65]: img.shape
```

```
Out[65]: (1140, 2000, 3)
```



```
In [85]: img[0:2,0:2,:]
```

```
Out[85]: array([[207, 213, 201],
               [210, 216, 204]],

            [[211, 217, 205],
             [216, 222, 210]]], dtype=uint8)
```

```
In [87]: img.flatten().shape
```

```
Out[87]: (6840000,)
```

- Now, again, the image is a point in space.
- But now the space is 6,840,000-dimensional!
- We'll talk about this towards the end of the course.

Data Cleaning

- ML+DM typically assume "clean" data.
- Ways that data might not be "clean":
 - noise (e.g., distortion on phone).
 - outliers (e.g., data entry or instrument error).
 - missing values (no value available or not applicable)
 - duplicated data (repetitions, or different storage formats).
- Any of these can lead to problems in analyses.
 - 1) want to fix these issues, apply data cleaning algorithms.
 - 2) some ML methods are robust to these.
 - often, ML is the best way to detect/fix these.

How much data do we need?

- A difficult if not impossible question to answer.
- Usual answer: "more is better".

- With the warning: "as long as the quality doesn't suffer".
- Another popular answer: "ten times the number of features".
 - I don't like this view. Features are not the enemy!

Feature aggregation

- Combine data to form new features
- Useful if there are few examples of a particular case

```
In [23]: titanic['deck'].value_counts()
```

```
Out[23]: C    59
         B    47
         D    33
         E    32
         A    15
         F    13
         G     4
         Name: deck, dtype: int64
```

```
In [93]: titanic_agg = titanic.copy()

# aggregate decks A and B into the "upper" deck category
titanic_agg["upper"] = titanic_agg['deck'].isin(("A","B"))
titanic_agg.tail()
```

```
Out[93]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	embark_town	alive	alone	upper
886	0	2	male	27.0	0	0	13.00	S	Second	man	True	NaN	Southampton	no	True	False
887	1	1	female	19.0	0	0	30.00	S	First	woman	False	B	Southampton	yes	True	True
888	0	3	female	NaN	1	2	23.45	S	Third	woman	False	NaN	Southampton	no	False	False
889	1	1	male	26.0	0	0	30.00	C	First	man	True	C	Cherbourg	yes	True	False
890	0	3	male	32.0	0	0	7.75	Q	Third	man	True	NaN	Queenstown	no	True	False

(Not shown: we should still fix up the NaNs here!)

Feature selection

```
In [94]: titanic_id = titanic.copy()

# Adding an irrelevant feature
titanic_id['id'] = titanic_id.index
titanic_id.head()
```

```
Out[94]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	embark_town	alive	alone	id
0	0	3	male	22.0	1	0	7.2500	S	Third	man	True	NaN	Southampton	no	False	0
1	1	1	female	38.0	1	0	71.2833	C	First	woman	False	C	Cherbourg	yes	False	1
2	1	3	female	26.0	0	0	7.9250	S	Third	woman	False	NaN	Southampton	yes	True	2
3	1	1	female	35.0	1	0	53.1000	S	First	woman	False	C	Southampton	yes	False	3
4	0	3	male	35.0	0	0	8.0500	S	Third	man	True	NaN	Southampton	no	True	4

- Remove features that are not relevant to the task.
- `id` probably not relevant for prediction.

Feature transformation

Discretization (binning): turn numerical data into categorical

```
In [95]: titanic['age'].head()
```

```
Out[95]: 0    22.0
1    38.0
2    26.0
3    35.0
4    35.0
Name: age, dtype: float64
```

```
In [108]: ages = pd.cut(titanic['age'], bins=(0,20,30,100))
ages.head()
```

```
Out[108]: 0      (20, 30]
1      (30, 100]
2      (20, 30]
3      (30, 100]
4      (30, 100]
Name: age, dtype: category
Categories (3, interval[int64, right]): [(0, 20] < (20, 30] < (30, 100]]
```

```
In [109]: ages_cat = pd.get_dummies(ages)
pd.concat([titanic['age'], ages_cat],axis=1).head()
```

```
Out[109]:
```

	age	(0, 20]	(20, 30]	(30, 100]
0	22.0	0	1	0
1	38.0	0	0	1
2	26.0	0	1	0
3	35.0	0	0	1
4	35.0	0	0	1

Mathematical transformations

- e.g. log, exp, square, sqrt, etc.
- also, scaling/normalization

```
In [115]: titanic.head()
```

```
Out[115]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	embark_town	alive	alone
0	0	3	male	22.0	1	0	7.2500	S	Third	man	True	NaN	Southampton	no	False
1	1	1	female	38.0	1	0	71.2833	C	First	woman	False	C	Cherbourg	yes	False
2	1	3	female	26.0	0	0	7.9250	S	Third	woman	False	NaN	Southampton	yes	True
3	1	1	female	35.0	1	0	53.1000	S	First	woman	False	C	Southampton	yes	False
4	0	3	male	35.0	0	0	8.0500	S	Third	man	True	NaN	Southampton	no	True

```
In [28]: titanic_mod = titanic.copy()

# fare --> sqrt(fare)
titanic_mod['fare'] = np.sqrt(titanic_mod['fare'])
titanic_mod.head()
```

```
Out[28]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	embark_town	alive	alone
0	0	3	male	22.0	1	0	2.692582	S	Third	man	True	NaN	Southampton	no	False
1	1	1	female	38.0	1	0	8.442944	C	First	woman	False	C	Cherbourg	yes	False
2	1	3	female	26.0	0	0	2.815138	S	Third	woman	False	NaN	Southampton	yes	True
3	1	1	female	35.0	1	0	7.286975	S	First	woman	False	C	Southampton	yes	False
4	0	3	male	35.0	0	0	2.837252	S	Third	man	True	NaN	Southampton	no	True

Example use case: something needs to be non-negative (exp) or shouldn't be non-negative (log).

Exploratory data analysis (EDA)

- You should always "look" at the data first.
- But how do you "look" at features and high-dimensional objects?
 - Summary statistics
 - Visualization

Categorical summary statistics

- Some summary statistics for a categorical variable:
 - **Frequencies** of different classes.
 - **Mode**: category that occurs most often.

```
In [120]: titanic['deck'].value_counts(normalize=True) # frequencies
```

```
Out[120]: C    0.290640  
         B    0.231527  
         D    0.162562  
         E    0.157635  
         A    0.073892  
         F    0.064039  
         G    0.019704  
         Name: deck, dtype: float64
```

```
In [121]: titanic['deck'].mode()[0]
```

```
Out[121]: 'C'
```

```
In [122]: titanic["survived"].mode()[0]
```

```
Out[122]: 0
```

```
In [123]: titanic.groupby("survived").size()
```

```
Out[123]: survived  
0      549  
1      342  
         dtype: int64
```

Continuous summary statistics

- Measures of location:

- **Mean:** average value.
- **Median:** value such that half points are larger/smaller.
- **Quantiles:** value such that t fraction of points are smaller.
- Measures of spread:
 - **Range:** minimum and maximum values.
 - **Variance:** measures how far values are from mean.
 - Square root of variance is **standard deviation**.
 - **Intequantile ranges:** difference between quantiles

In [146]: `titanic.describe()`

Out[146]:

	survived	pclass	age	sibsp	parch	fare
count	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

In [137]: `titanic['fare'].var()`

Out[137]: 2469.436845743117

Notice that the mean and std are sensitive to extreme values:

```
In [152]: data = [0,1,2,3,3,5,7,8,9,10,14,15,17,1000000] # the "1000000" is an outlier
print("Mean with outlier  :", np.mean(data))
print("Mean without outlier:", np.mean(data[:-1]))
```

```
Mean with outlier  : 71435.28571428571
Mean without outlier: 7.230769230769231
```

```
In [153]: print("Std with outlier  :", np.std(data))
print("Std without outlier:", np.std(data[:-1]))
```

```
Std with outlier  : 257537.51466268493
Std without outlier: 5.351546809515718
```

Whereas the median is not:

```
In [154]: print("Median with outlier  :", np.median(data))
print("Median without outlier:", np.median(data[:-1]))
```

```
Median with outlier  : 7.5
Median without outlier: 7.0
```

Distances and similarities

- There are also summary statistics between features.
 - Hamming distance:
 - Number of elements in the vectors that aren't equal.
 - Euclidean distance:
 - How far apart are the vectors?
 - Correlation:
 - Does one increase/decrease linearly as the other increases?
 - Between -1 and 1.

Limitations of summary statistics

- Summary statistics can be misleading
- A famous example is [Anscombe's quartet](https://en.wikipedia.org/wiki/Anscombe%27s_quartet) (https://en.wikipedia.org/wiki/Anscombe%27s_quartet), four datasets with:
 - Almost same means.
 - Almost same variances.
 - Almost same correlations.
 - Almost same linear fits.
 - Look completely different.

In [155]: *# Code below from seaborn documentation: https://seaborn.pydata.org/examples/anscombes_quartet.html*

Load the example dataset for Anscombe's quartet

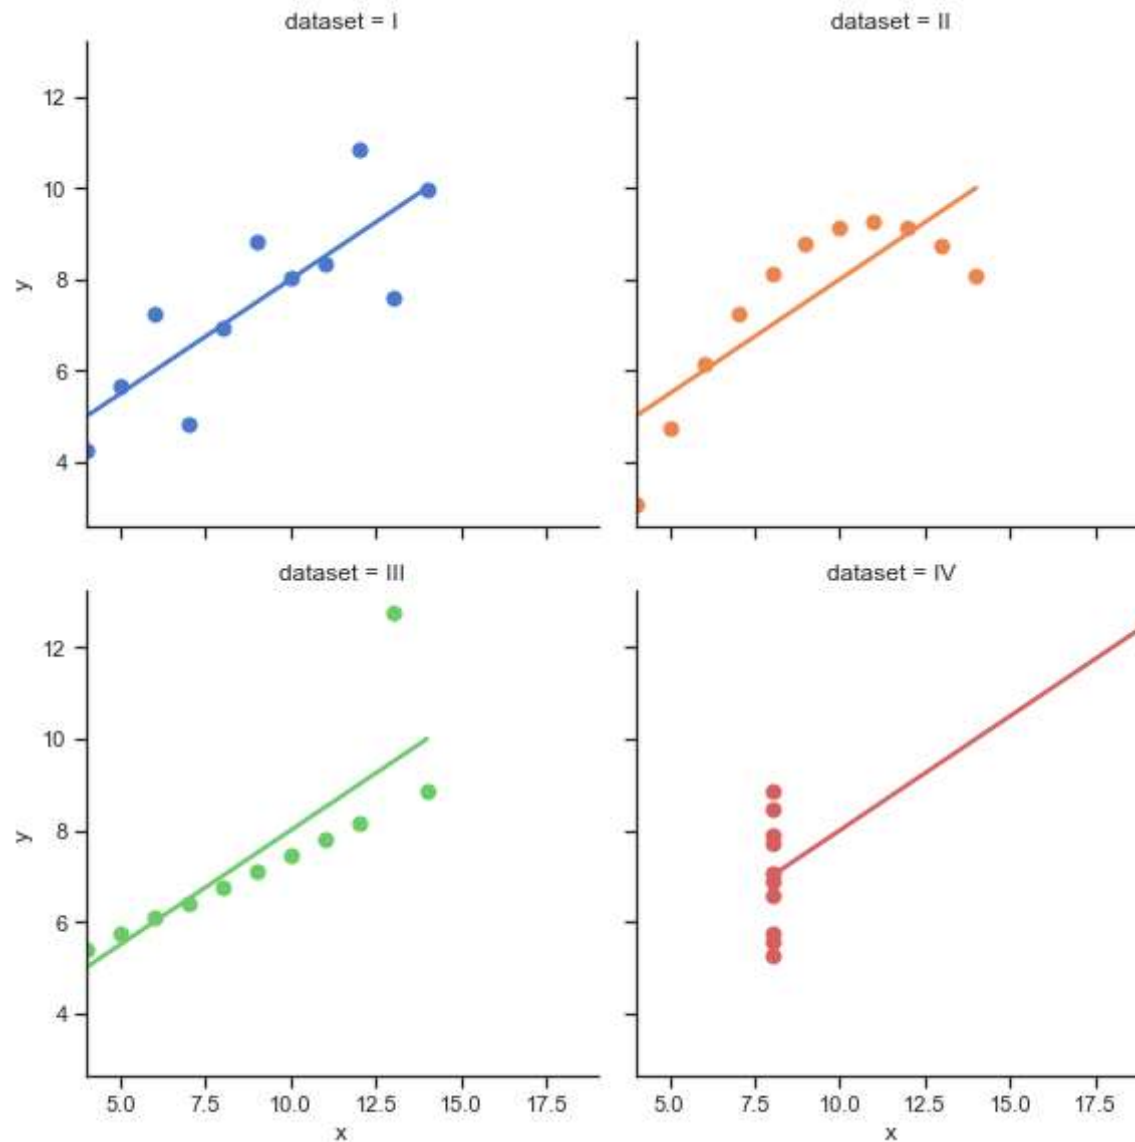
```
anscombe = sns.load_dataset("anscombe")
```

Show the results of a linear regression within each dataset

```
sns.lmplot(x="x", y="y", col="dataset", hue="dataset", data=anscombe,  
           col_wrap=2, ci=None, palette="muted", size=4,  
           scatter_kws={"s": 50, "alpha": 1});
```

C:\Users\User\anaconda3\lib\site-packages\seaborn\regression.py:580: UserWarning: The `size` parameter has been renamed to `height`; please update your code.

```
warnings.warn(msg, UserWarning)
```




Visualization

- You can learn a lot from 2D plots of the data:
 - Patterns, trends, outliers, unusual patterns.

- We'll use the `matplotlib` library to do most of our basic plotting.
- For fancier plots, you can try `seaborn`.


Basic plot

```
In [178]: x = np.linspace(0,10,100)  
plt.plot(x, np.cos(x));
```



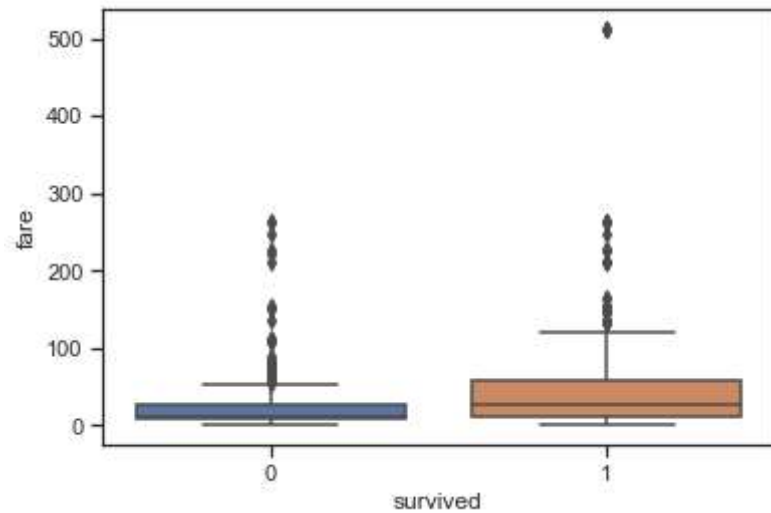
Histogram

```
In [190]: plt.hist(titanic['age'])  
plt.xlabel('age')  
plt.ylabel('frequency');  
# sns.distplot(iris["sepal_length"]);
```




Box plot

```
In [195]: sns.boxplot(x="survived", y="fare", data=titanic);
```




```
In [202]: sns.boxplot(x="survived", y="age", data=titanic);
```



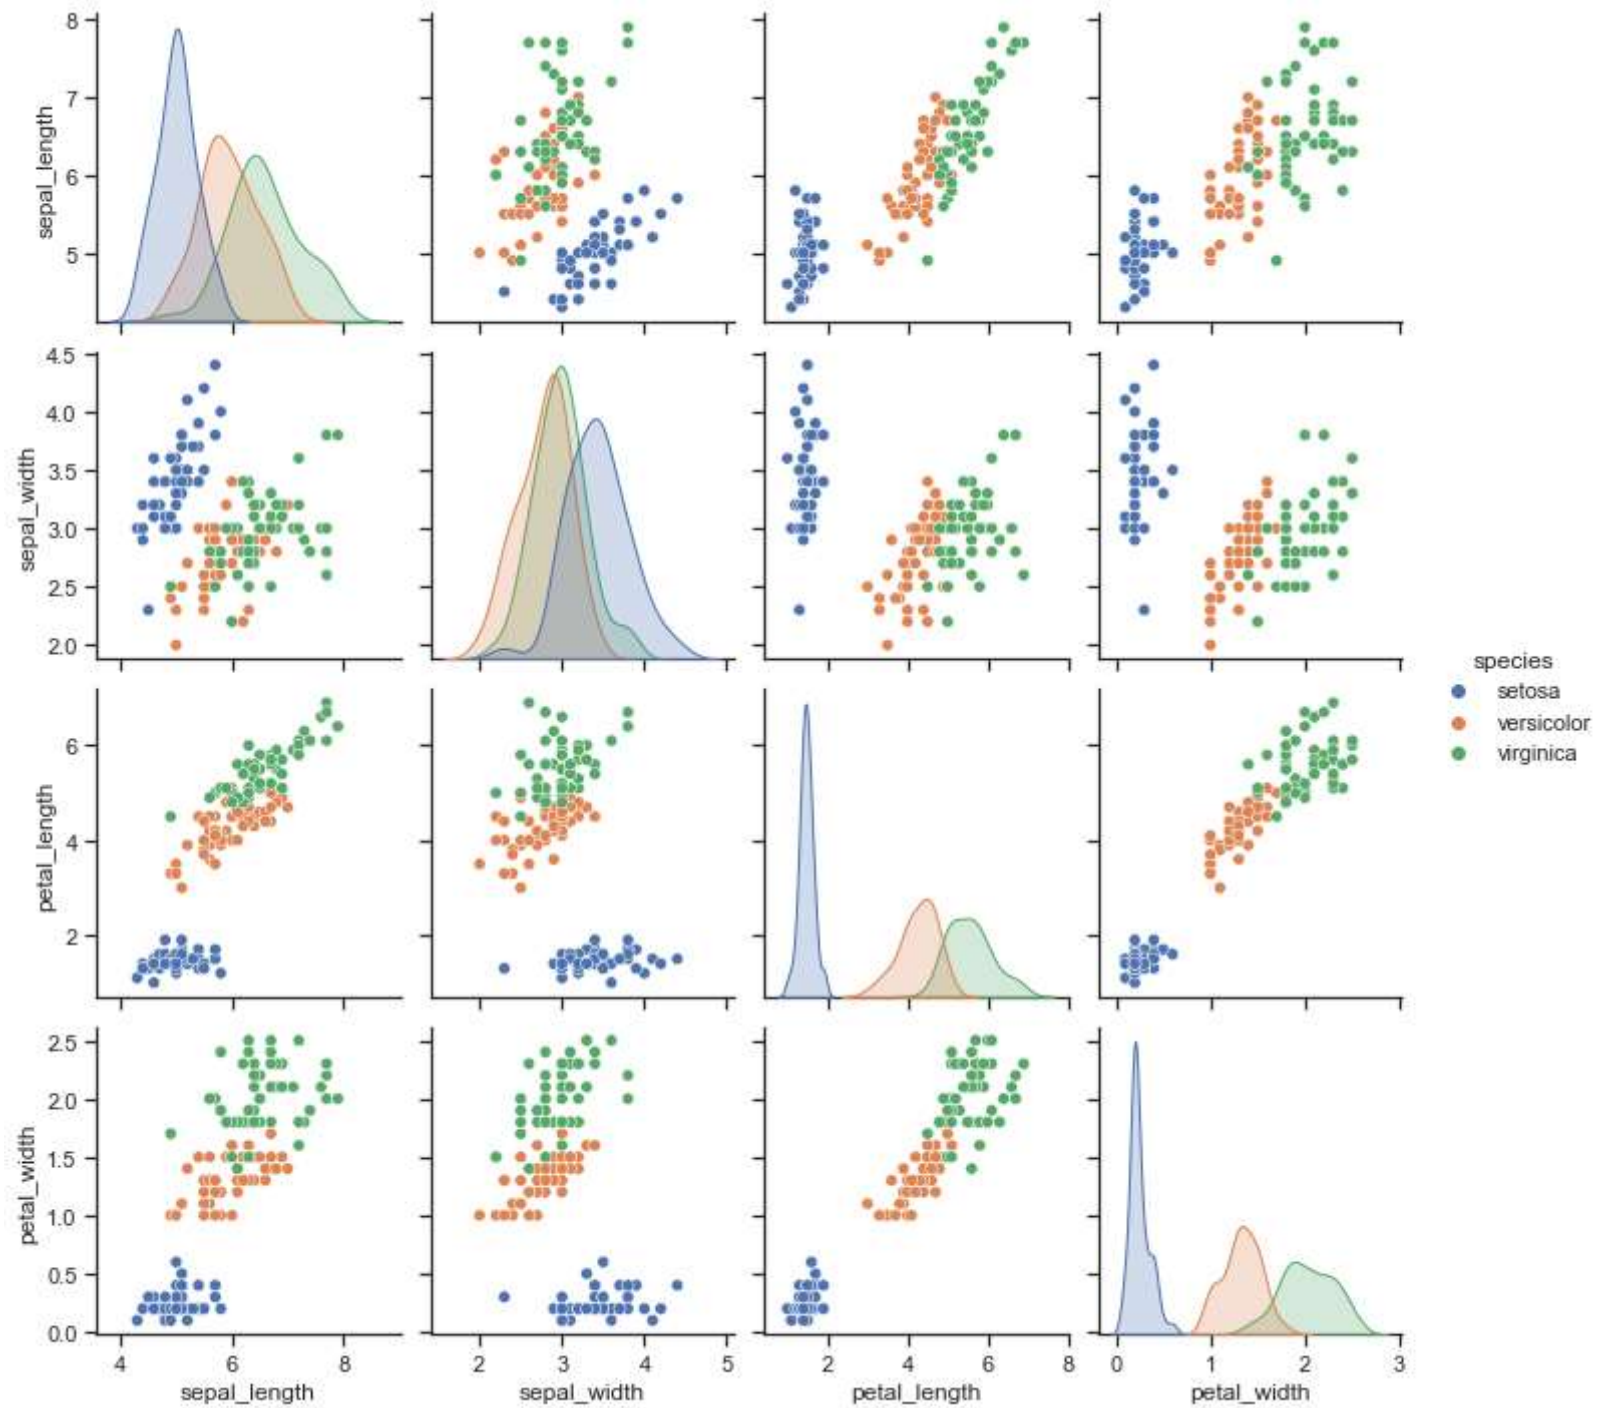
Scatterplot

```
In [208]: plt.scatter(iris['sepal_length'], iris['petal_length'])  
plt.xlabel('sepal_length')  
plt.ylabel('petal_length');
```



Scatterplot array

```
In [206]: sns.pairplot(iris, hue="species");
```

Summary

- Typical data mining steps:
 - Involves data collection, preprocessing, analysis, and evaluation.
- Object-feature representation and categorical/numerical features.
 - Transforming non-vector objects to vector representations.
- Feature transformations:
 - To address coupon collecting or simplify relationships between variables.
- Exploring data:
 - Summary statistics and data visualization.
- Post-lecture bonus slides: other visualization methods.