

---

# **Embedded Systems**

---

## **Real-Time Scheduling**

**Dr. Anas A. Youssef**  
**[anas.youssef@ci.menofia.edu.eg](mailto:anas.youssef@ci.menofia.edu.eg)**

*Operating  
Systems:  
Internals  
and  
Design  
Principles*

# Chapter 10 Multiprocessor, Multicore and Real-Time Scheduling

Eighth Edition  
By William Stallings

# Real-Time Systems

- The operating system, and in particular the scheduler, is perhaps the most important component

## Examples:

- control of laboratory experiments
- process control in industrial plants
- robotics
- air traffic control
- telecommunications
- military command and control systems



- Correctness of the system depends not only on **the logical result** of the computation but also on **the time at which the results are produced**
- Tasks or processes attempt to control or react to events that take place in the outside world
- These events occur in “real time” and tasks must be able to keep up with them



# Hard and Soft Real-Time Tasks

## Hard real-time task

- one that **must** meet its deadline
- otherwise it will cause unacceptable damage or a fatal error to the system

## Soft real-time task

- has an associated deadline that is **desirable** but not mandatory
- it still makes sense to schedule and complete the task even if it has passed its deadline



# Periodic and Aperiodic Tasks

## ■ Periodic tasks

- requirement may be stated as:
  - once per period  $T$
  - exactly  $T$  units apart

## ■ Aperiodic tasks

- has a deadline by which it must finish or start
- may have a constraint on both start and finish time

# Characteristics of Real Time Systems

Real-time operating systems have requirements in five general areas:

Determinism

Responsiveness

User control

Reliability

Fail-soft operation



# Determinism

- Concerned with how long an operating system delays before acknowledging an interrupt
- Operations are performed at fixed, predetermined times or within predetermined time intervals
  - when multiple processes are competing for resources and processor time, no system will be fully deterministic

The extent to which an operating system can deterministically satisfy requests depends on:

the speed with which it can respond to interrupts

whether the system has sufficient capacity to handle all requests within the required time

# Responsiveness


- Together with determinism make up the response time to external events
  - critical for real-time systems that must meet timing requirements imposed by individuals, devices, and data flows external to the system
- Concerned with how long, after acknowledgment, it takes an operating system to service the interrupt

## Responsiveness includes:

- amount of time required to initially handle the interrupt and begin execution of the interrupt service routine (ISR)
- amount of time required to perform the ISR
- effect of interrupt nesting



# User Control



- Generally much broader in a real-time operating system than in ordinary operating systems
- It is essential to allow the user fine-grained control over task priority
- User should be able to distinguish between hard and soft tasks and to specify relative priorities within each class
- May allow user to specify such characteristics as:

paging or process swapping    what processes must always be resident in main memory    what disk transfer algorithms are to be used    what rights the processes in various priority bands have

# Reliability

- More important for real-time systems than non-real time systems
- Real-time systems respond to and control events in real time so loss or degradation of performance may have catastrophic consequences such as:
  - financial loss
  - major equipment damage
  - loss of life



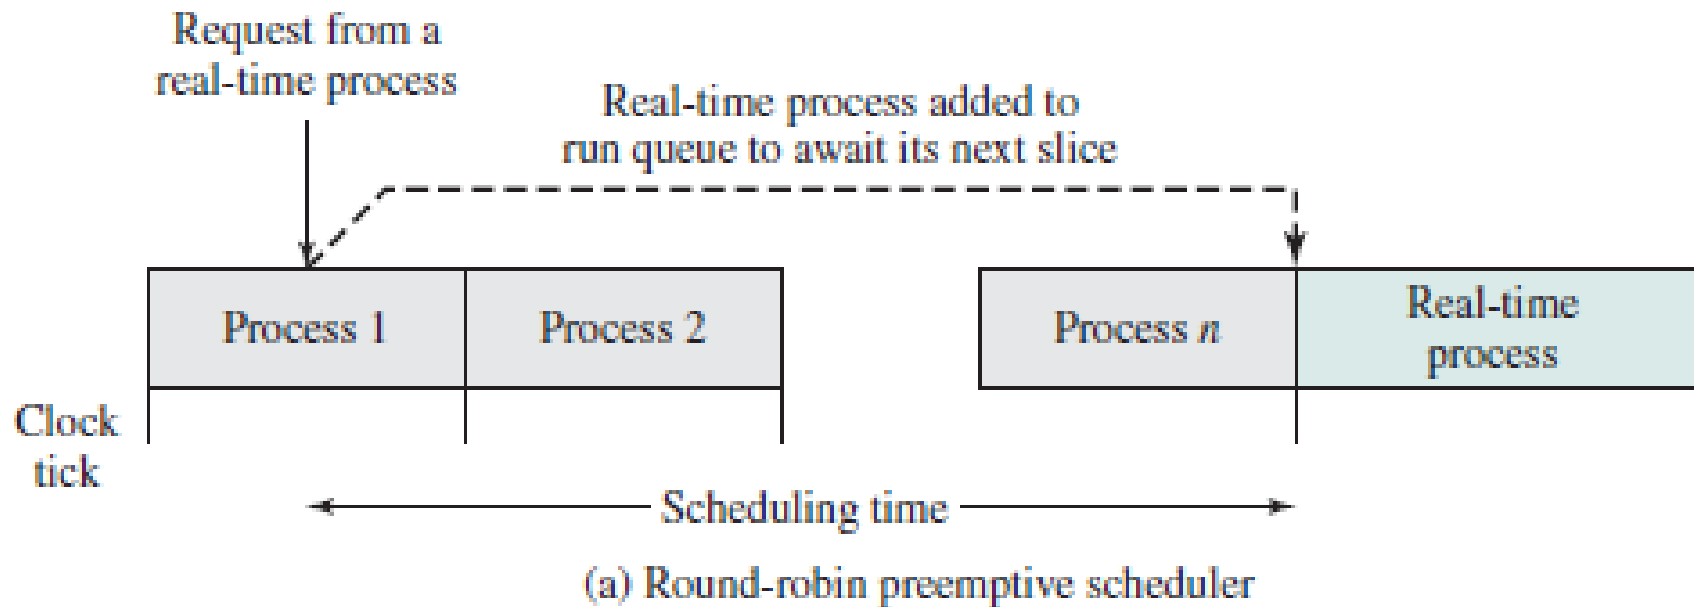
# Fail-Soft Operation

- A characteristic that refers to the ability of a system to fail in such a way as to preserve as much capability and data as possible
- Important aspect is stability
  - a real-time system is stable if the system will meet the deadlines of its most critical, highest-priority tasks even if some less critical task deadlines are not always met

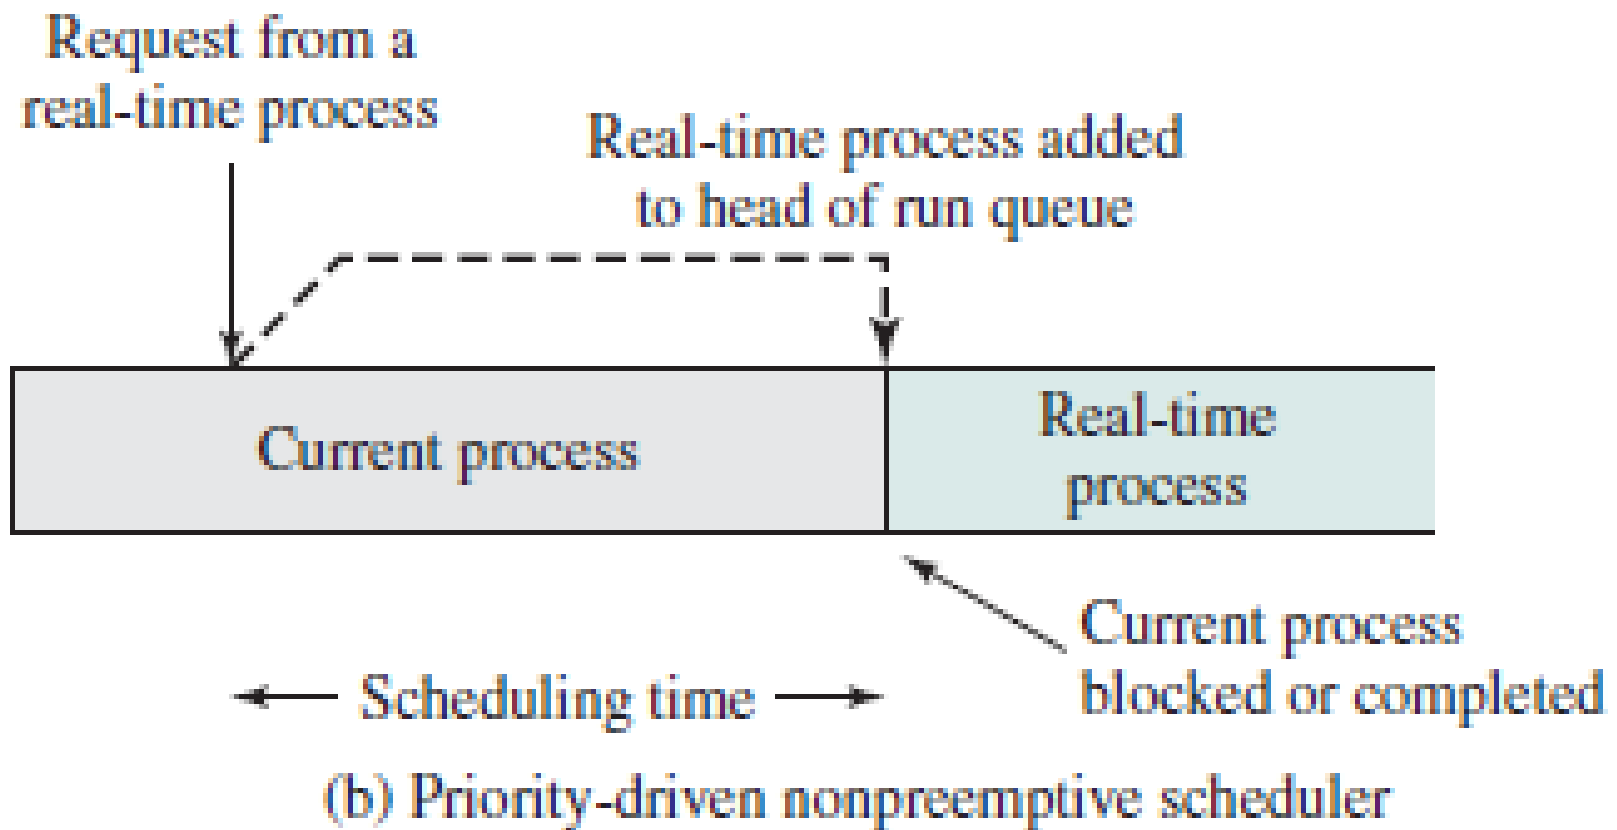




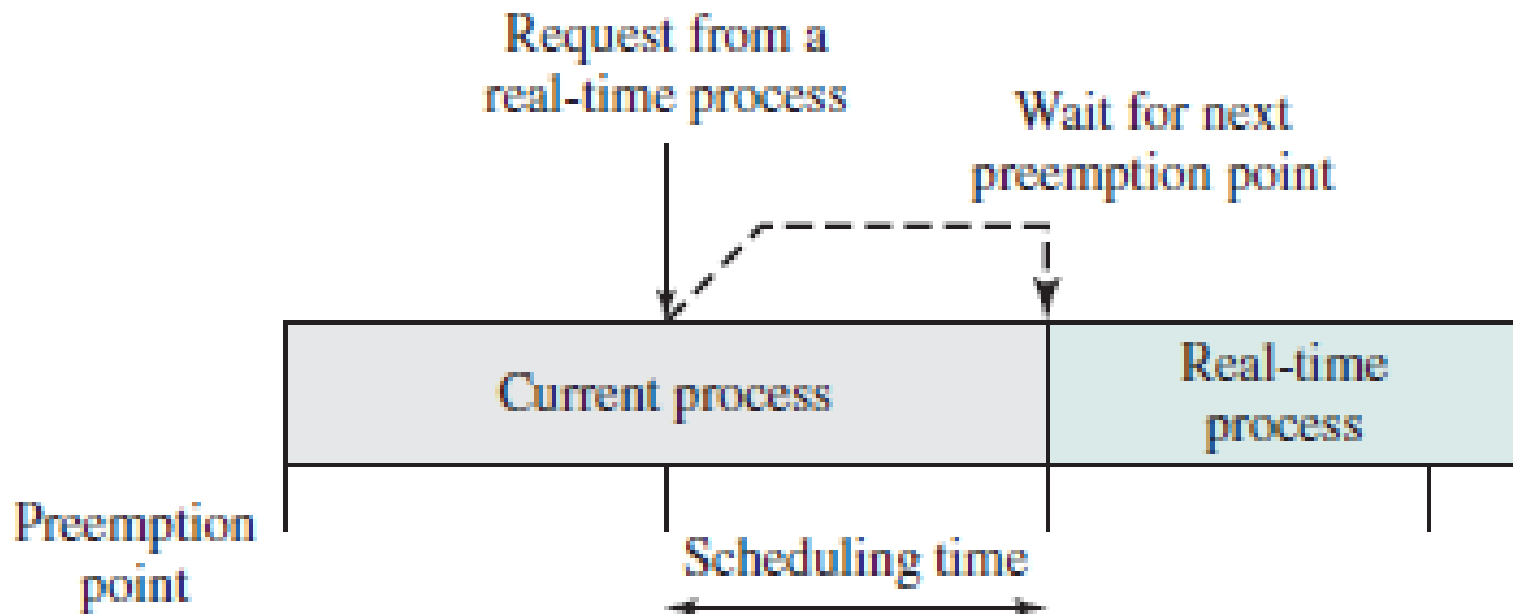
# Scheduling of Real-Time Process



# Scheduling of Real-Time Process



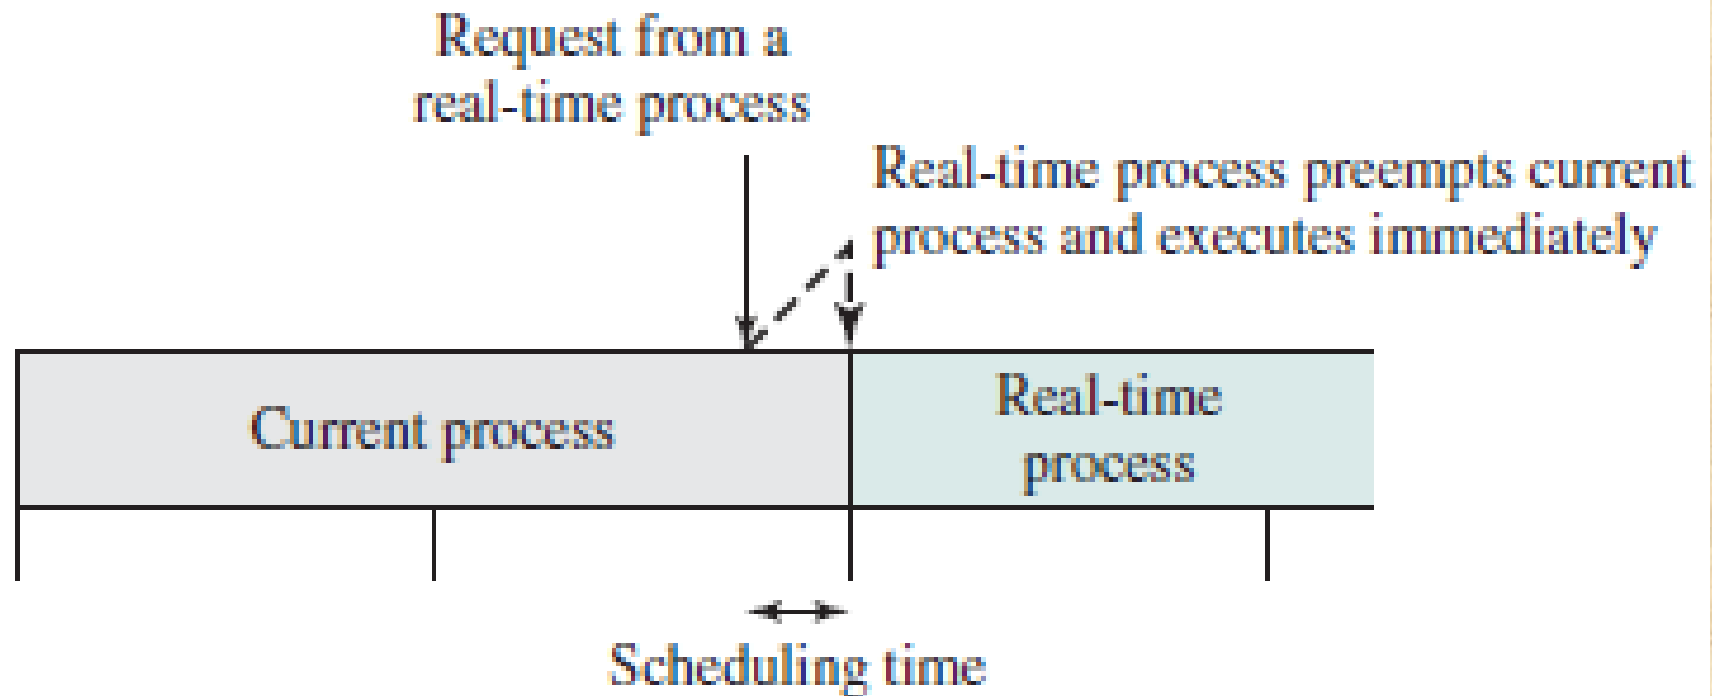
# Scheduling of Real-Time Process



(c) Priority-driven preemptive scheduler on preemption points


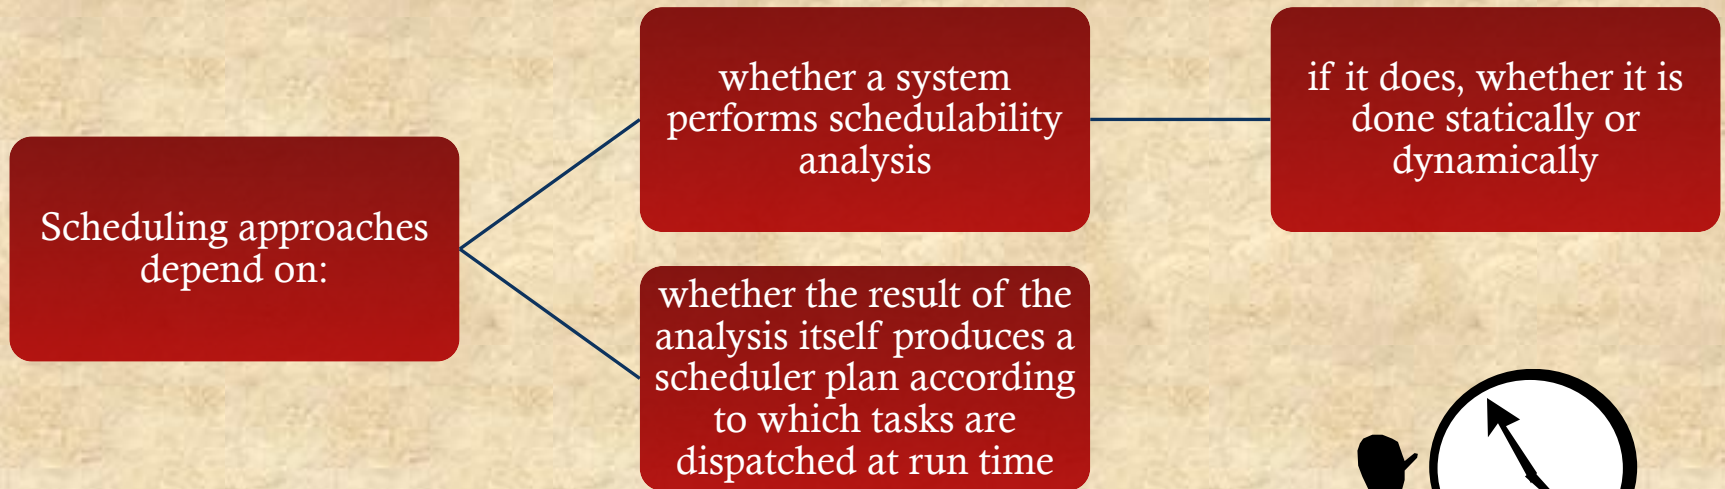


# Scheduling of Real-Time Process



(d) Immediate preemptive scheduler

# Real-Time Scheduling



# Classes of Real-Time Scheduling Algorithms

**Static table-driven**

**Static priority-driven preemptive**

**Dynamic planning-based**

**Dynamic best effort**



# Static table-driven

- performs a static analysis of feasible schedules of dispatching which result in a schedule that determines, at run time, when a task must begin execution
- applicable to tasks that are **periodic**
- Input to the analysis consists of the periodic arrival time, execution time, periodic ending deadline, and relative priority of each task.
- The scheduler **attempts to develop a schedule** that enables it to meet the requirements of all periodic tasks.
- This is a predictable approach but one that is **inflexible**, because any change to any task requirements requires that the schedule be redone.
- Example: **Earliest-deadline-first**

# Static priority-driven preemptive

- a static analysis is performed but no schedule is drawn up
- **analysis is used to assign priorities to tasks** so that a traditional priority-driven preemptive scheduler can be used
- One example of this approach is the **rate monotonic scheduling** algorithm , which assigns static priorities to tasks based on the length of their periods.



# Dynamic planning-based

- feasibility is determined at run time (dynamically) rather than offline prior to the start of execution (statically).
- an arriving task is accepted for execution only if it is feasible to meet its time constraints
- one result of the analysis is a schedule or plan that is used to decide when to dispatch this task
- an attempt is made to create a schedule that contains the previously scheduled tasks as well as the new arrival.
- If the new arrival can be scheduled in such a way that its deadlines are satisfied and that no currently scheduled task misses a deadline, then the schedule is revised to accommodate the new task.



# Dynamic best effort

- no feasibility analysis is performed
- when a task arrives, the system assigns a priority based on the characteristics of the task
- system tries to meet all deadlines and aborts any started process whose deadline is missed
- Typically, the tasks are aperiodic and so no static scheduling analysis is possible.
- With this type of scheduling, until a deadline arrives or until the task completes, we do not know whether a timing constraint will be met. This is the major disadvantage of this form of scheduling.
- Its advantage is that it is easy to implement.

# Deadline Scheduling

- Real-time operating systems are designed with the objective of starting real-time tasks as rapidly as possible and emphasize rapid interrupt handling and task dispatching
- Real-time applications are generally not concerned with sheer speed but rather with completing (or starting) tasks at the most valuable times
- Priorities provide a crude tool and do not capture the requirement of completion (or initiation) at the most valuable time



# Information Used for Deadline Scheduling

## Ready time

- time task becomes ready for execution

## Resource requirements

- resources required by the task while it is executing

## Starting deadline

- time task must begin

## Priority

- measures relative importance of the task

## Completion deadline

- time task must be completed

## Processing time

- time required to execute the task to completion

## Subtask scheduler

- a task may be decomposed into a mandatory subtask and an optional subtask



# Table 10.3

## Execution Profile of Two Periodic Tasks

Process	Arrival Time	Execution Time	Ending Deadline
A(1)	0	10	20
A(2)	20	10	40
A(3)	40	10	60
A(4)	60	10	80
A(5)	80	10	100
•	•	•	•
•	•	•	•
•	•	•	•
B(1)	0	25	50
B(2)	50	25	100
•	•	•	•
•	•	•	•
•	•	•	•

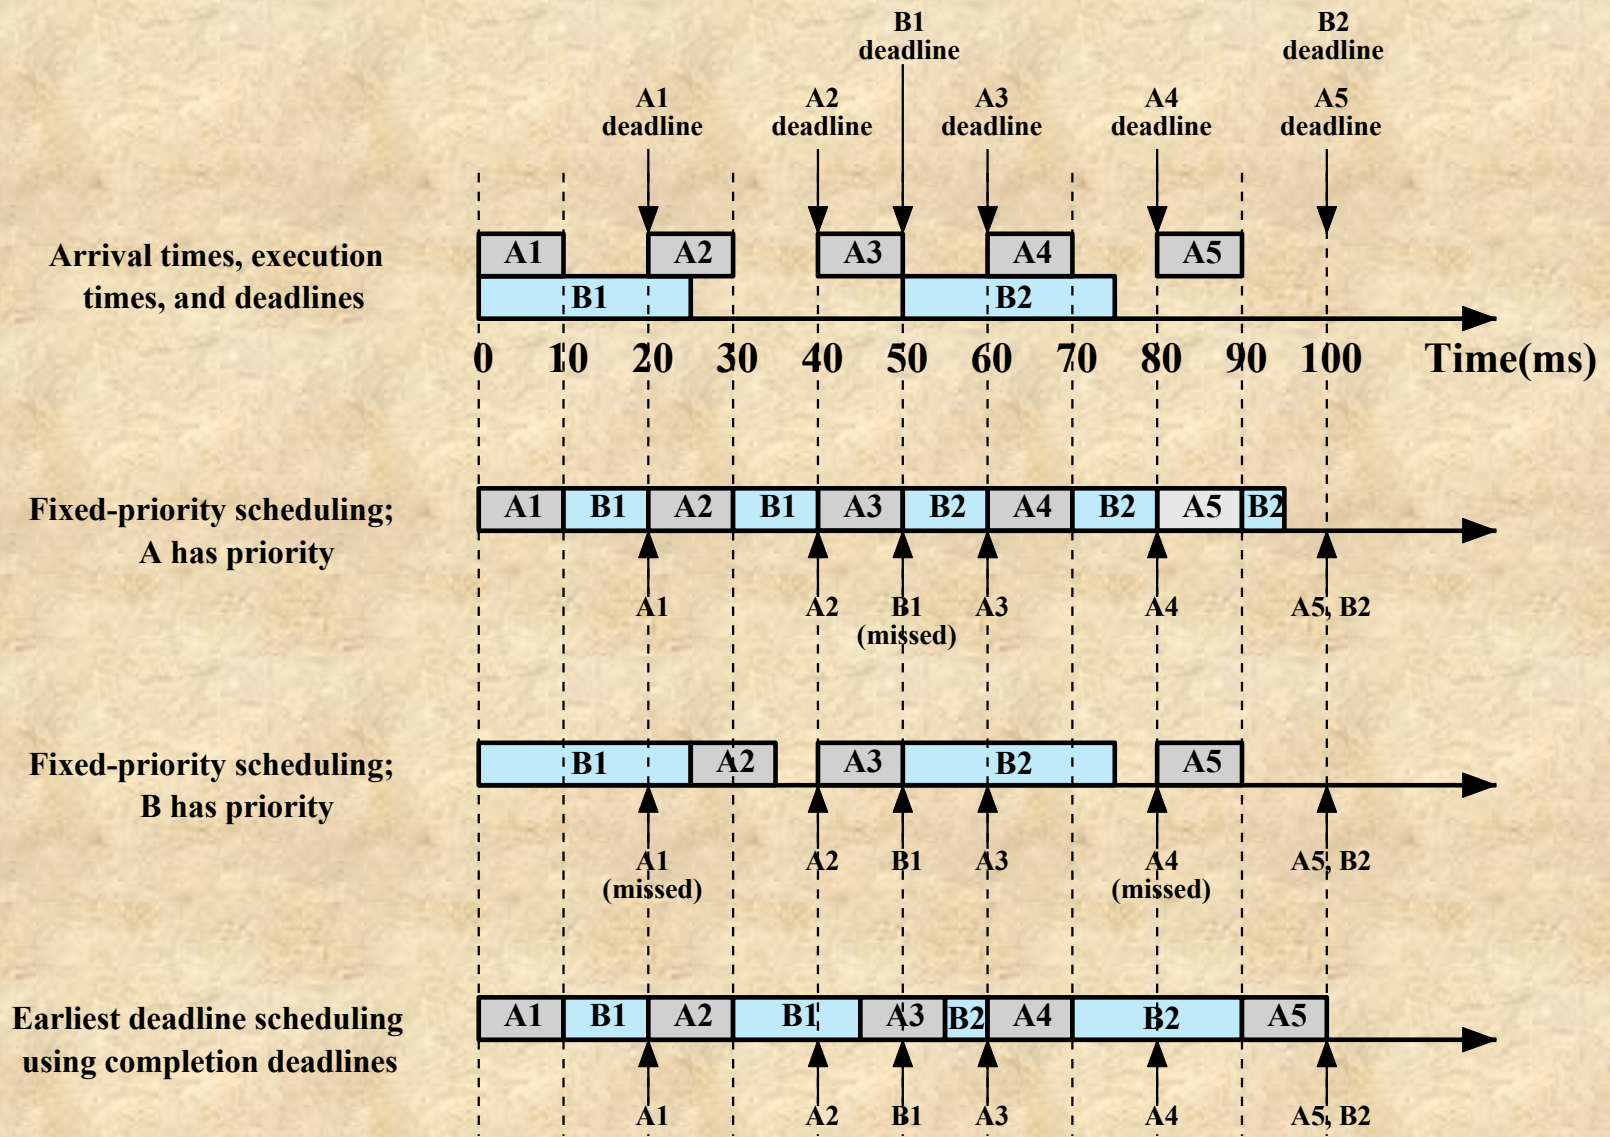


Figure 10.5 Scheduling of Periodic Real-time Tasks with Completion Deadlines (based on Table 1 10.3)

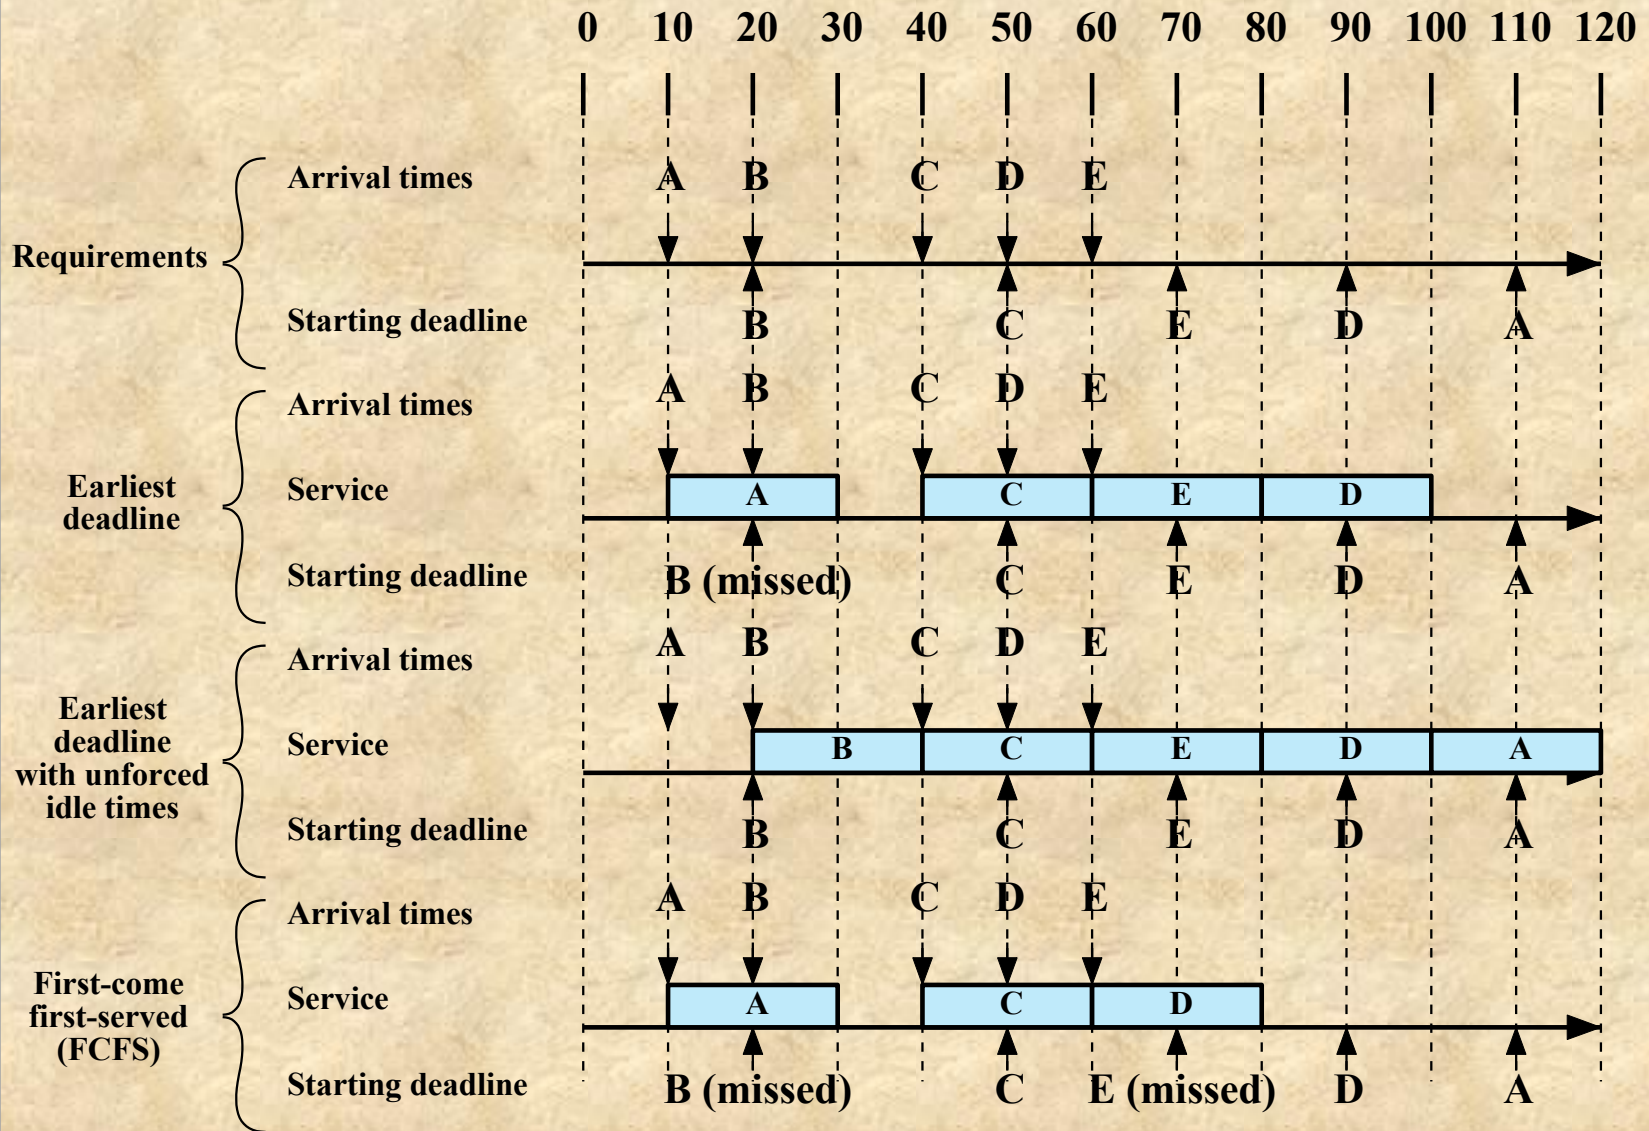
---

## Table 10.4

### Execution Profile of Five Aperiodic Tasks

Process	Arrival Time	Execution Time	Starting Deadline
A	10	20	110
B	20	20	20
C	40	20	50
D	50	20	90
E	60	20	70



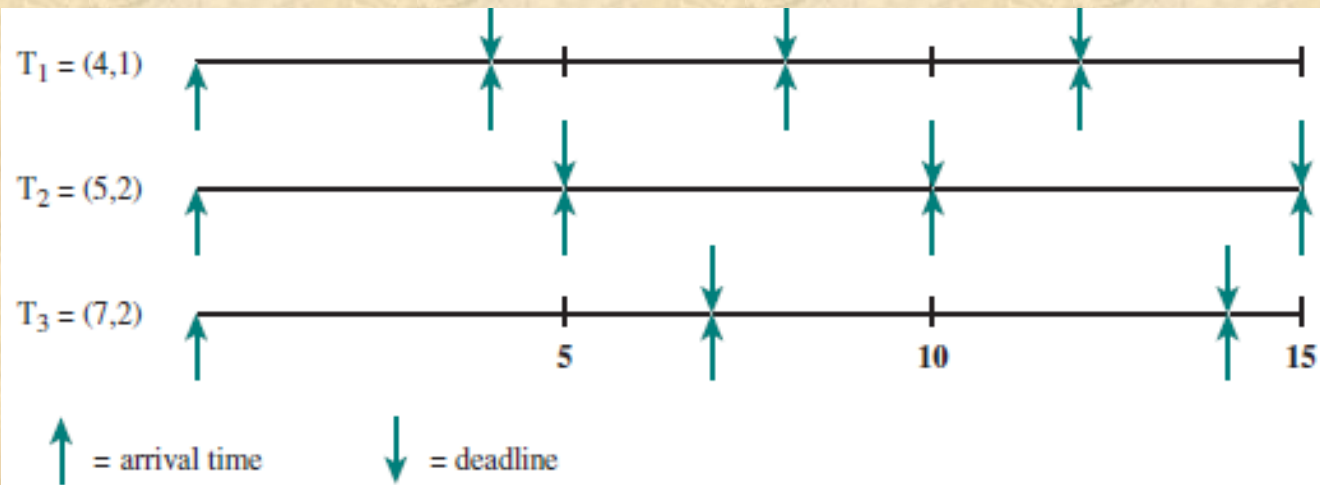


**Figure 10.6 Scheduling of Aperiodic Real-time Tasks with Starting Deadlines**

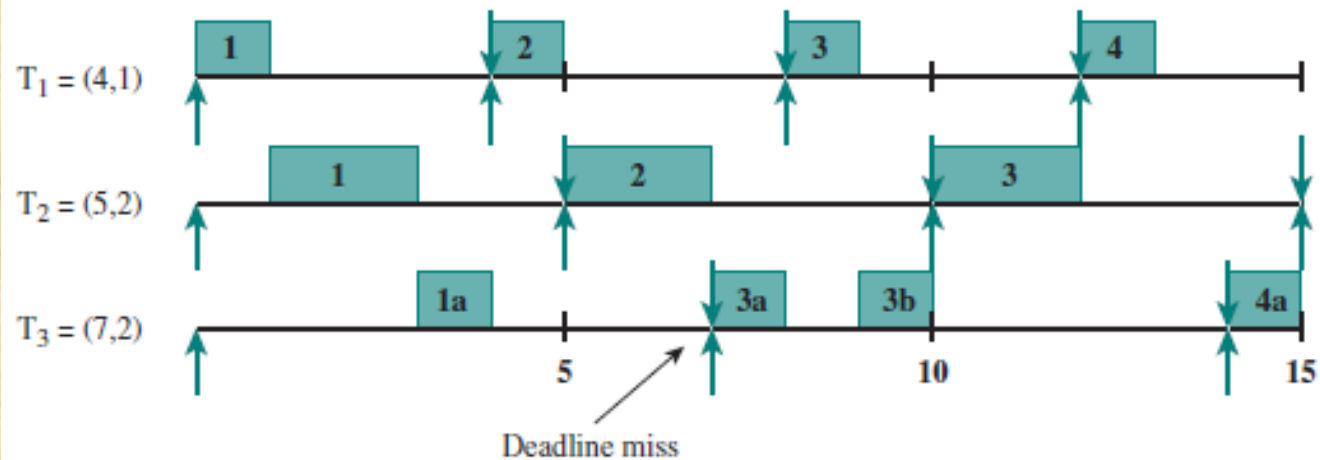


## Rate Monotonic Scheduling (RMS)

Figure 10.7 A Task Set with RMS



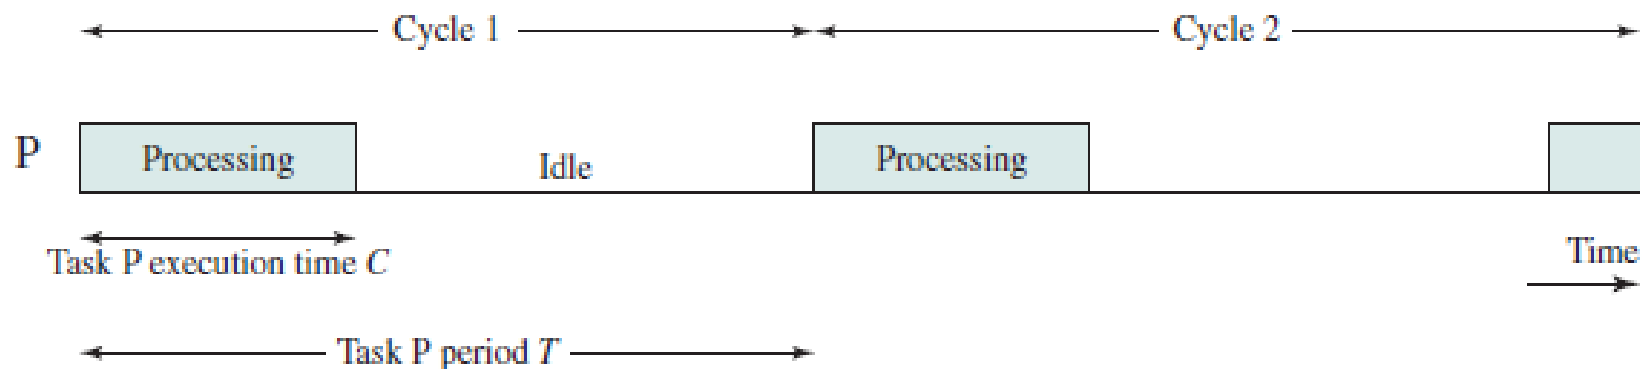
(a) Arrival times and deadlines for task  $T_i = (P_i, C_i)$ ;  
 $P_i$  = period,  $C_i$  = processing time



(b) Scheduling results

## Rate Monotonic Scheduling (RMS) Example





**Figure 10.8** Periodic Task Timing Diagram

## For Perfect Scheduling

$$\frac{C_1}{T_1} + \frac{C_2}{T_2} + \dots + \frac{C_n}{T_n} \leq 1$$


## For RMS only

$$\frac{C_1}{T_1} + \frac{C_2}{T_2} + \dots + \frac{C_n}{T_n} \leq n(2^{1/n} - 1)$$

**Table 10.5**

**Value of  
the RMS  
Upper  
Bound**

$n$	$n(2^{1/n} - 1)$
1	1.0
2	0.828
3	0.779
4	0.756
5	0.743
6	0.734
•	•
•	•
•	•
¥	$\ln 2 \gg 0.693$



As an example, consider the case of three periodic tasks, where  $U_i = C_i/T_i$ :

- **Task P<sub>1</sub>:**  $C_1 = 20$ ;  $T_1 = 100$ ;  $U_1 = 0.2$
- **Task P<sub>2</sub>:**  $C_2 = 40$ ;  $T_2 = 150$ ;  $U_2 = 0.267$
- **Task P<sub>3</sub>:**  $C_3 = 100$ ;  $T_3 = 350$ ;  $U_3 = 0.286$

The total utilization of these three tasks is  $0.2 + 0.267 + 0.286 = 0.753$ . The upper bound for the schedulability of these three tasks using RMS is

$$\frac{C_1}{T_1} + \frac{C_2}{T_2} + \frac{C_3}{T_3} \leq n(2^{1/3} - 1) = 0.779$$




# Priority Inversion

- Can occur in any priority-based preemptive scheduling scheme
- Particularly relevant in the context of real-time scheduling
- Best-known instance involved the Mars Pathfinder mission in 1997
- Occurs when circumstances within the system force a higher priority task to wait for a lower priority task
- If a lower-priority task has locked a resource and a higher-priority task attempts to lock that resource, the higher-priority task will be put in a blocked state until the resource is available.
- If the lower-priority task soon finishes with the resource and releases it, the higher-priority task may quickly resume and it is possible that no real-time constraints are violated.

# Unbounded Priority Inversion

- The duration of a priority inversion depends not only on the time required to handle a shared resource, but also on the unpredictable actions of other unrelated tasks

# Unbounded Priority Inversion



(a) Unbounded priority inversion



# Unbounded Priority Inversion

$t_1$ :  $T_3$  begins executing.

$t_2$ :  $T_3$  locks semaphore  $s$  and enters its critical section.

$t_3$ :  $T_1$ , which has a higher priority than  $T_3$ , preempts  $T_3$  and begins executing.

$t_4$ :  $T_1$  attempts to enter its critical section but is blocked because the semaphore is locked by  $T_3$ ;  $T_3$  resumes execution in its critical section.

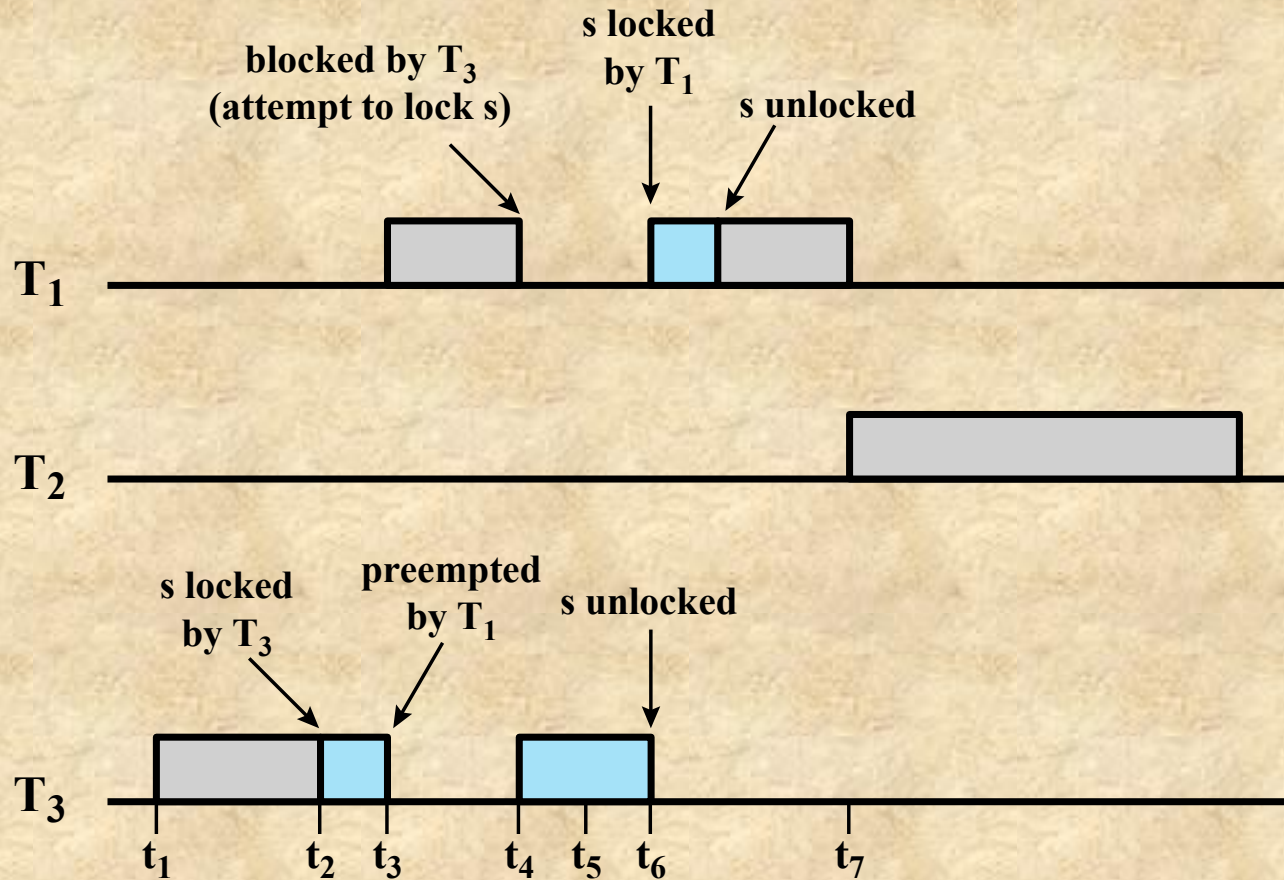
$t_5$ :  $T_2$ , which has a higher priority than  $T_3$ , preempts  $T_3$  and begins executing.

$t_6$ :  $T_2$  is suspended for some reason unrelated to  $T_1$  and  $T_3$ ;  $T_3$  resumes.

$t_7$ :  $T_3$  leaves its critical section and unlocks the semaphore.  $T_1$  preempts  $T_3$ , locks the semaphore, and enters its critical section.

- $T_1$  must wait for both  $T_3$  and  $T_2$  to complete and fails.

# Priority Inheritance



(b) Use of priority inheritance

# Priority Inheritance

$t_1$ :  $T_3$  begins executing.

$t_2$ :  $T_3$  locks semaphore  $s$  and enters its critical section.

$t_3$ :  $T_1$ , which has a higher priority than  $T_3$ , preempts  $T_3$  and begins executing.

$t_4$ :  $T_1$  attempts to enter its critical section but is blocked because the semaphore is locked by  $T_3$ .  $T_3$  is immediately and temporarily assigned the same priority as  $T_1$ .  $T_3$  resumes execution in its critical section.

$t_5$ :  $T_2$  is ready to execute but, because  $T_3$  now has a higher priority,  $T_2$  is unable to preempt  $T_3$ .

$t_6$ :  $T_3$  leaves its critical section and unlocks the semaphore: its priority level is downgraded to its previous default level.  $T_1$  preempts  $T_3$ , locks the semaphore, and enters its critical section.

$t_7$ :  $T_1$  is suspended for some reason unrelated to  $T_2$ , and  $T_2$  begins executing.



# Priority Ceiling

- A priority is associated with each resource.
- The priority assigned to a resource is one level higher than the priority of its highest priority user.
- The scheduler then dynamically assigns this priority to any task that accesses the resource.
- Once the task finishes with the resource, its priority returns to normal.

# Summary

- Real-time scheduling
  - Background
  - Characteristics of real-time operating systems
  - Real-time scheduling
  - Deadline scheduling
  - Rate monotonic scheduling
  - Priority inversion