

Patterns

- Observer Pattern

Accessing Hardware Patterns

- Hardware Proxy Pattern
- Hardware Adapter Pattern
- Mediator Pattern
- Interrupt Pattern
- Polling Pattern

Concurrency and Resource Management

- Cyclic Executive Pattern
- Critical Region Pattern

Lecture 6

- **Pattern Hatching** – Locating the Right Patterns
- **Pattern Mining** – Rolling Your Own Patterns
- **Pattern Instantiation** – Applying Patterns in Your Designs
- **Design pattern** is a generalized solution to a commonly occurring problem.

The solution must be **general** enough to be applied in a **wide set** of application domains.

● A **good design** is composed of a set of design patterns applied to a piece of functional software, achieves a balanced optimization of the design criteria, incurs an acceptable cost

- **Basic Structure of a Design Pattern**

- **Name**

- provides a “**handle**” or means to reference the pattern.

- **Purpose**

- provides the problem context and the QoS aspects the pattern seeks to optimize.

- specifies under which situations the pattern is appropriate and under which situations it should be avoided

- **Solution**

- structure and behavior of the pattern.
 - elements of the pattern and their roles in the pattern context.

- **Consequences**

- set of pros and cons of the use of the pattern.

Pattern Hatching

Familiarize yourself with patterns -> Identify Design Criteria -> Rank Design Criteria by Criticality -> Identify Potential Design Alternatives -> Evaluate Design Alternatives -> Select and Apply Design Alternatives -> Verify Design Functionality, if functionality broken, go back to identify design alternatives. Else -> Verify Design Goals are met, if not, go back to identify design alternatives, else -> End.

Some Common Design Optimization Criteria

- Performance

- Worst case
- Average case

- Predictability

- Schedulability

- Throughput: number of tasks to be executed per second

- Average: Average time to execute task
- Sustained: not working average, but always working at peak potential (ثبات الأداء مع زيادة التاسكات)
- Burst: القدرة على التعامل مع فترات الضغط المفاجئة زي وقت نزول النتيجة

- Reliability

- With respect to errors or failures

Design Tradeoff Spreadsheet

We list Design criteria and their weights in rows and Design solution in columns, Filling cells of each criteria and solution with a score, total weighted score for every design solution is calculated by: **Sum of Design criteria weight * Design solution score. The highest Design Solution's total weighted score is the answer**

Design Tradeoff Spreadsheet

Table 2-3: Design tradeoff spreadsheet.

Design Solution	Design Criteria					Total Weighted Score
	Criteria 1	Criteria 2	Criteria 3	Criteria 4	Criteria 5	
	Weight = 7	Weight = 5	Weight = 3	Weight = 2	Weight = 1.5	
	Score	Score	Score	Score	Score	
Alternative 1	7	3	6	9	4	106
Alternative 2	4	8	5	3	4	95
Alternative 3	10	2	4	8	8	120
Alternative 4	2	4	9	7	6	84

Table 2-4: Design tradeoffs for ECG monitor system.

Design Solution	Design Criteria				Total Weighted Score
	Efficiency	Maintainability	Flexibility	Memory Usage	
	Weight = 7	Weight = 5	Weight = 4	Weight = 7	
	Score	Score	Score	Score	
Client Server	3	7	8	5	123
Push	8	4	7	9	167
Observer	8	7	9	9	190

Observer Design Pattern

Provides a means for objects to “listen in” on others while requiring **no modifications** to the data servers.

From Embedded Perspective, sensor data can be easily shared to other elements.

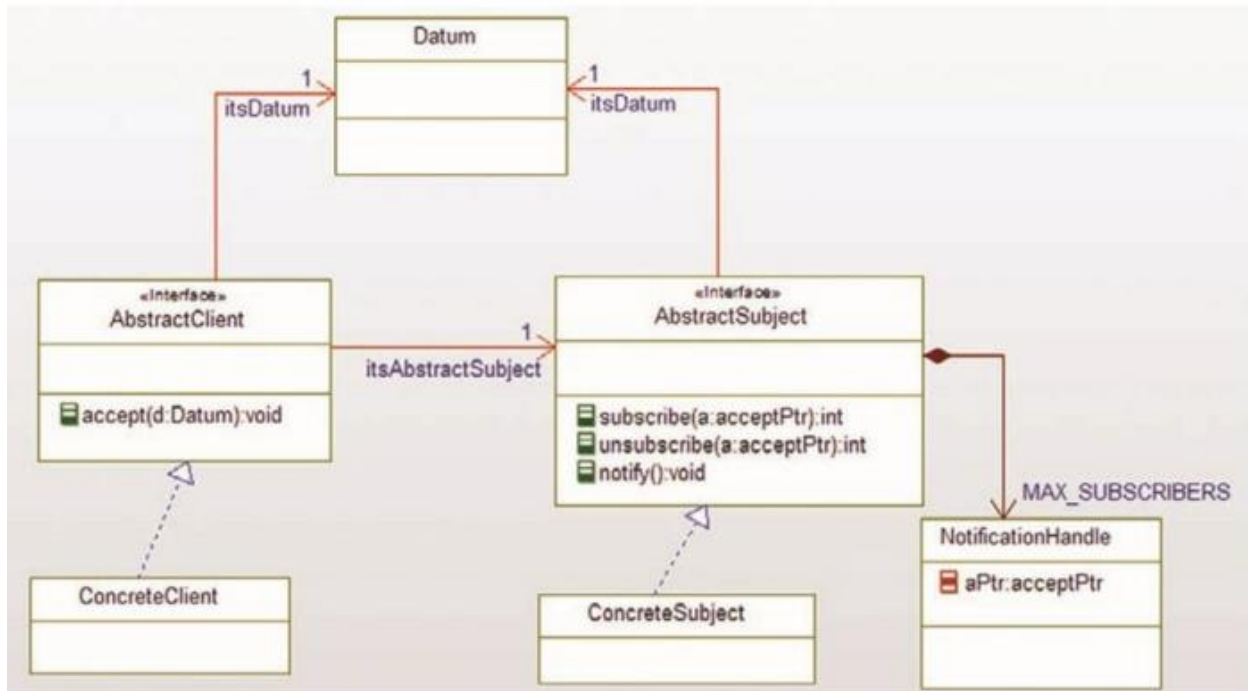
Abstract

- known as the “Publish-Subscribe Pattern”
- Provides **notification** to a set of interested clients that **relevant data have changed**.
- It does this **without requiring** the data server to have **any a priori knowledge** about its clients.
- Clients(Sensors) can use Subscribe function to add themselves to the notification list.
- The data server can then enforce whatever notification policy it desires.

Problem

- Each client can **request data periodically** from a data server in case the **data have changed**.
- If the data server **pushes** the data out, then it **must know who all of its clients** are.
- **subscription** and **un-subscription** services to data servers are allowed to clients.
- The pattern allows **dynamic modification of subscriber lists**.
- The server can enforce the appropriate **update policy to the notification** of its interested clients.

Structure



- Abstract Client Interface: has an accept function to accept data from the datum
- Abstract Subject Interface: The server, which has two functions to subscribe and unsubscribe
- Datum: the data shared among clients
- Notification Handle: Class handles the notification list and subscribers (add or remove subscribers)

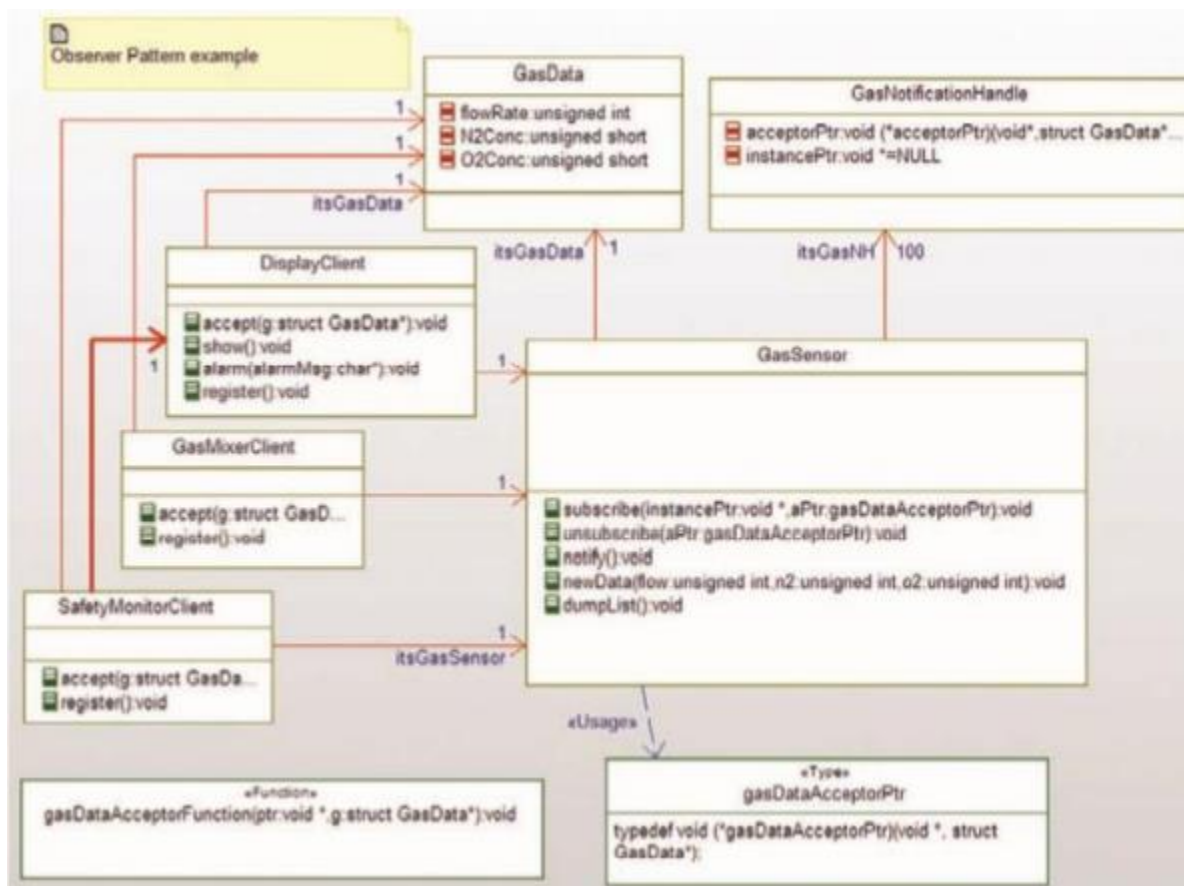
Consequences

- **Simplifies** the process of **distributing data** to a set of clients.
- **Maintains** the fundamental client-server relation while providing run-time **flexibility** of adding clients.
- **Compute efficiency is maintained** since clients can only know updates when data is changed.

Implementation Complexities

- The most complex aspects of this pattern are the **implementation of the notification handle and the management of the notification handle list**.
- The **easiest** approach for the notification list is to **declare an array big enough** to hold all potential clients. **This wastes memory** in highly dynamic systems with many potential clients.
- Another approach is to construct a **linked list** of all clients.

Example



Lecture 7

- Categories of Design Patterns include:

- Accessing hardware
- Concurrency and Resource Management
- State Machine implementation and usage
- Safety and Reliability

- Hardware Proxy Pattern (الوسيط)

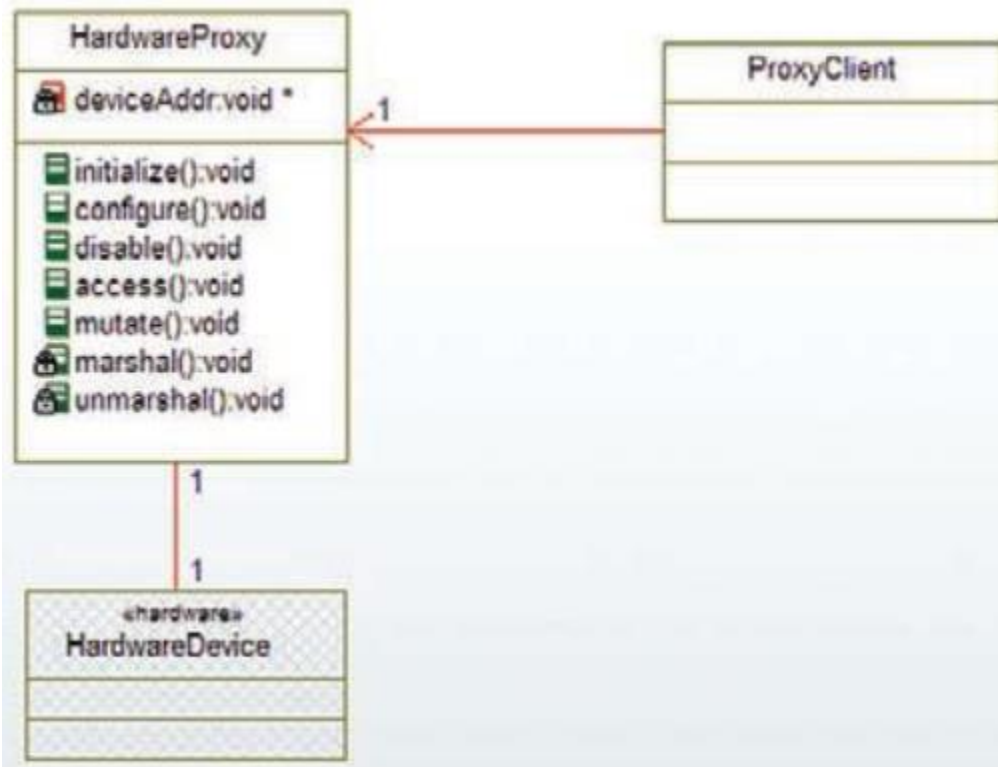
Abstract

- Creates a software element responsible for access to a piece of hardware and encapsulation of hardware compression and coding implementation.
- Uses a class (or struct) to encapsulate all access to a hardware device, regardless of its physical interface.
- Provides an encoding and connection-independent interface for clients and so promotes easy modification should the nature of the device interface or connection change

Problem

- By providing a proxy to sit between the clients and the actual hardware, the impact of hardware changes is greatly limited, easing any modifications in the hardware.

Structure



- `marshal()` – converts presentation (client) data format to native (motor) format
- `unmarshal()` – converts native (motor) data format to presentation (client) format

Consequences

- It provides **flexibility** for the actual hardware interface to **change** with absolutely **no changes in the clients**.
- This means that the **proxy clients are usually unaware of the native format of the data** and manipulate them only in presentation format.
- This can, however, have a **negative impact on runtime performance**.
- It may be more **time efficient** for the clients to **know the encoding details** and manipulate the data in **their native format**.

● Hardware Adapter Pattern (Like a convertor)

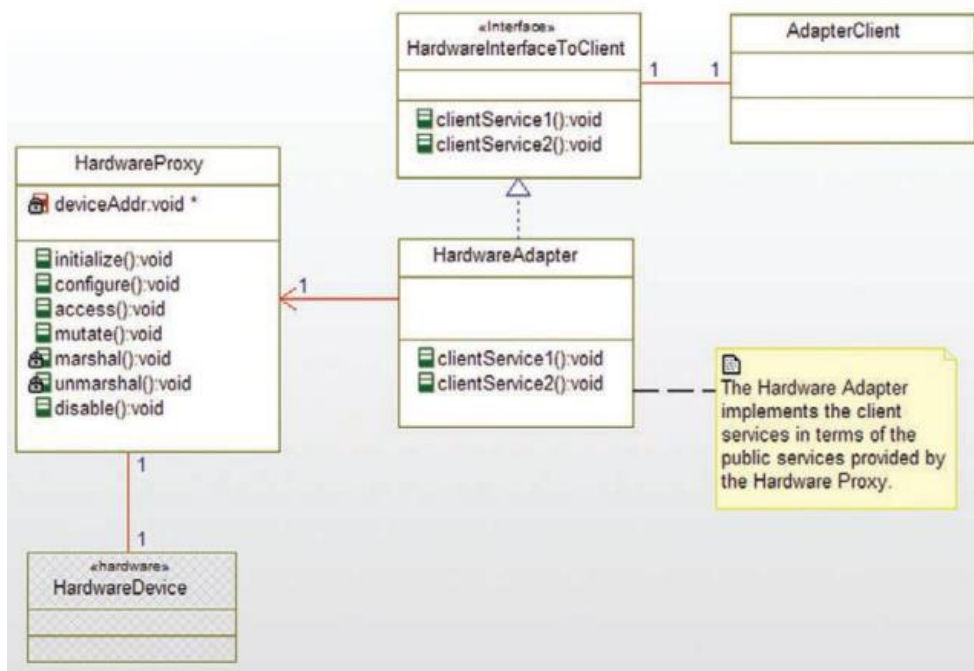
Abstract

- Provides a way of **adapting an existing hardware interface** into the expectations of the application.
- It is **useful** when the application requires or **uses one interface**, but the actual hardware provides another.
- Creates an element that **converts** between the two interfaces.

Problem

- It is useful when you have hardware that meets the semantic need of the system but that **has an incompatible interface**.
- The goal of the pattern is to **minimize the reworking** of code when one hardware design or implementation is replaced with another.

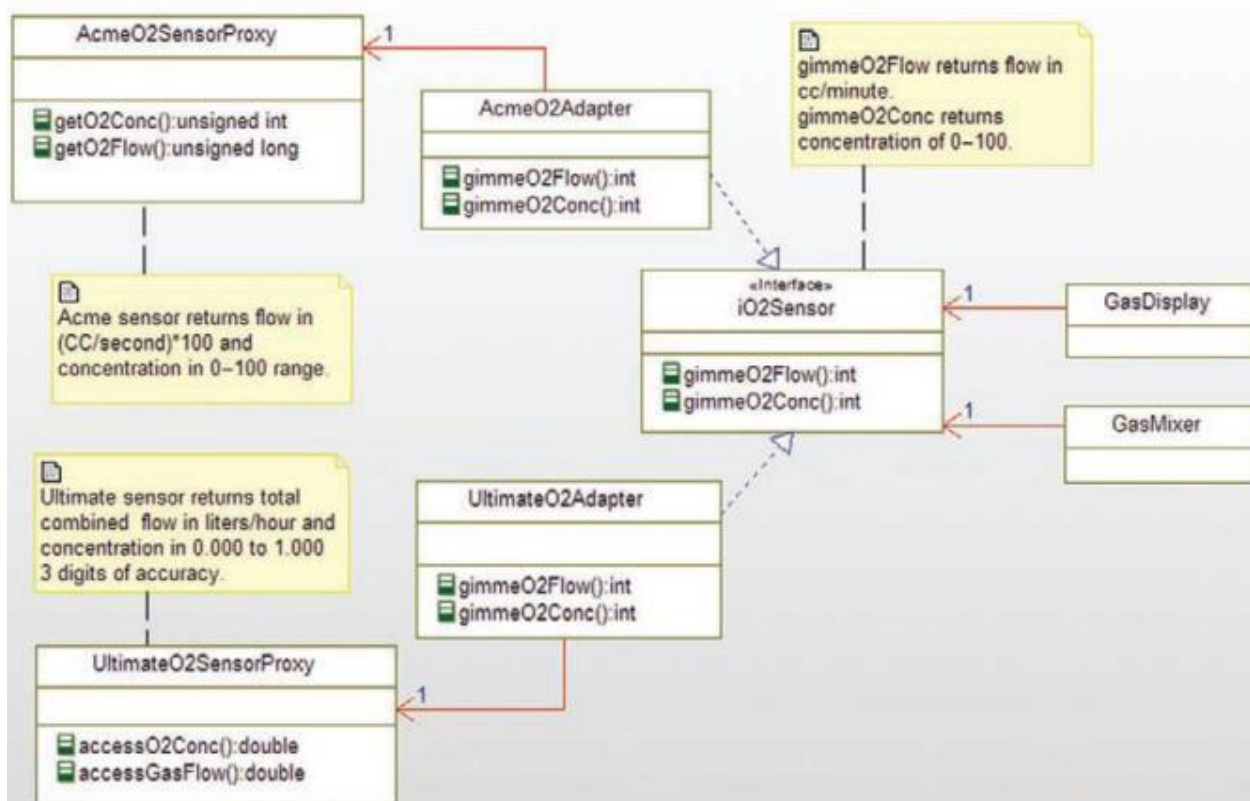
Structure



Consequences

- It allows various Hardware Proxies and their related hardware devices to be used as-is in different applications, while at the same time allowing existing applications to use different hardware devices **without change**.
- However, this pattern adds a level of indirection and therefore decreases runtime performance.

Example



Lecture 8

● Mediator Pattern (شبه ظابط المرور)

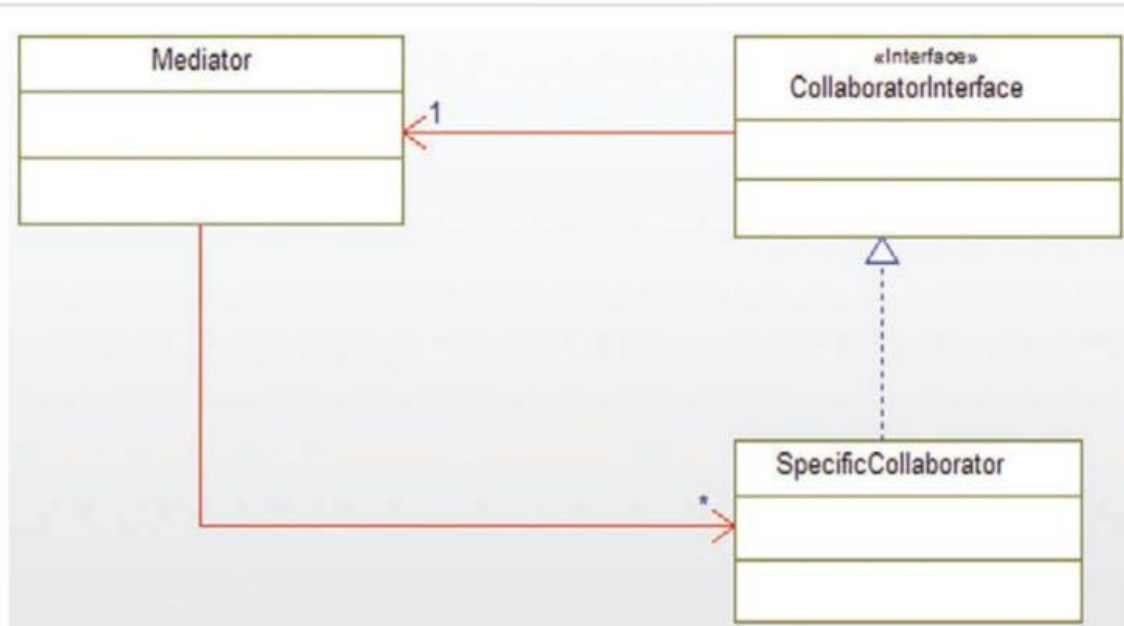
Abstract

- Provides a means of **coordinating a complex interaction** among a set of elements.
- Useful for **managing different hardware** elements when their behavior must be coordinated in well-defined but complex ways.

Problem

- Many embedded applications control **sets of actuators that must work in concert** to achieve the desired effect.
- For example, to achieve a coordinated movement of a multi-joint robot arm, all the motors must work together to provide the desired arm movement.

Structure



- It uses a mediator class to coordinate the actions of a set of collaborating devices to achieve the desired overall effect.
- It coordinates the control of the set of multiple Specific Collaborators (their number is indicated by the '*' multiplicity on the association between the Mediator and the Specific Collaborator)
- The Collaborator Interface is a specification of a set of services common to all Specific Collaborators that may be invoked by the Mediator.
- For example, it is common to have reset(), shutdown, initialize() style operations for all the hardware devices to facilitate bringing them all to a known initial, recovery and/or shutdown state.
- The Specific Collaborator represents one device and so may be a device driver or Hardware Proxy.
- It receives command messages from the Mediator and also sends messages to the Mediator when events of interest occur.

Consequences

- This pattern creates a mediator that coordinates the set of collaborating actuators but without requiring direct coupling of those devices.
- This greatly simplifies the overall design by minimizing the points of coupling and encapsulating the coordination within a single element.
- Whenever the Collaborator would have directly contacted another Collaborator, instead it notifies the Mediator who can decide how to respond as a collective collaborative whole.
- Since many embedded systems must react with high time fidelity, delays between the actions may result in unstable or undesirable effects.
- It is important that the mediator class can react within those time constraints.

● Interrupt Pattern

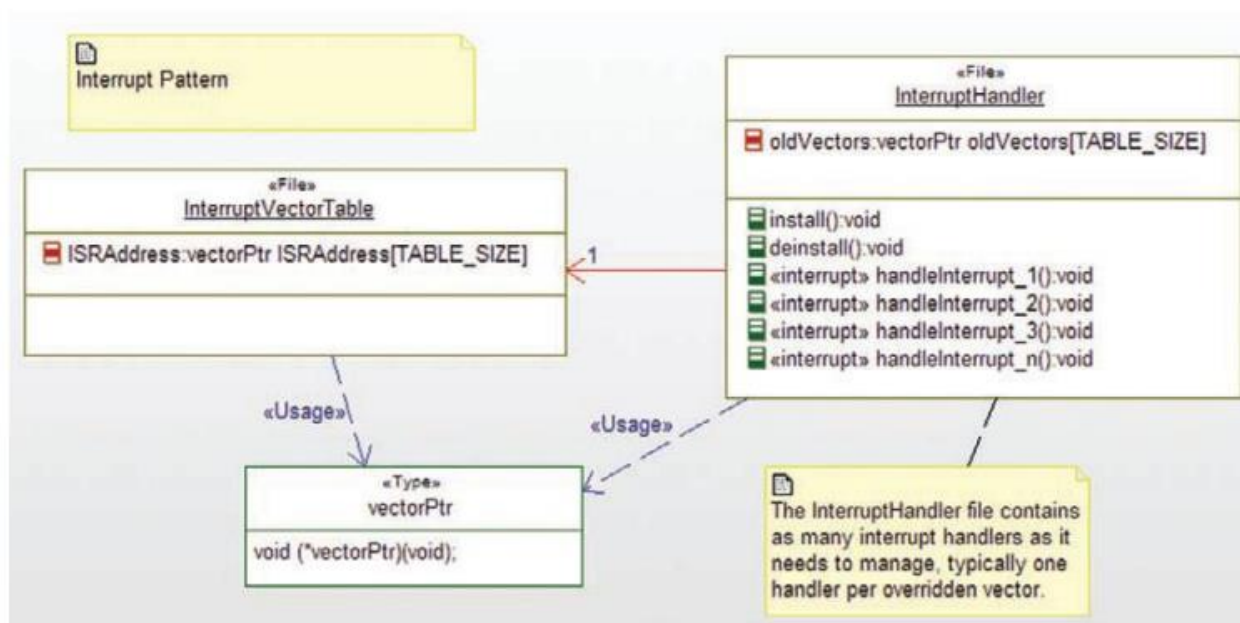
Abstract

- The Interrupt Pattern is a way of structuring the system to **respond** appropriately to incoming **events**.
- Once initialized, an interrupt will **pause its normal processing, handle the incoming event, and then return** the system to its original computations.

Problem

- In many systems, events have **different levels of urgency**.
- Most embedded systems have at least some events with a high urgency that must be handled even when the system is busy doing other processing.

Structure



Interrupt Vector Table: a table contains interrupt service routines addresses of interrupt handlers

Interrupt Handler: a file contains all interrupts

VectorPtr: a pointer to point an interrupt handler to its address in the table

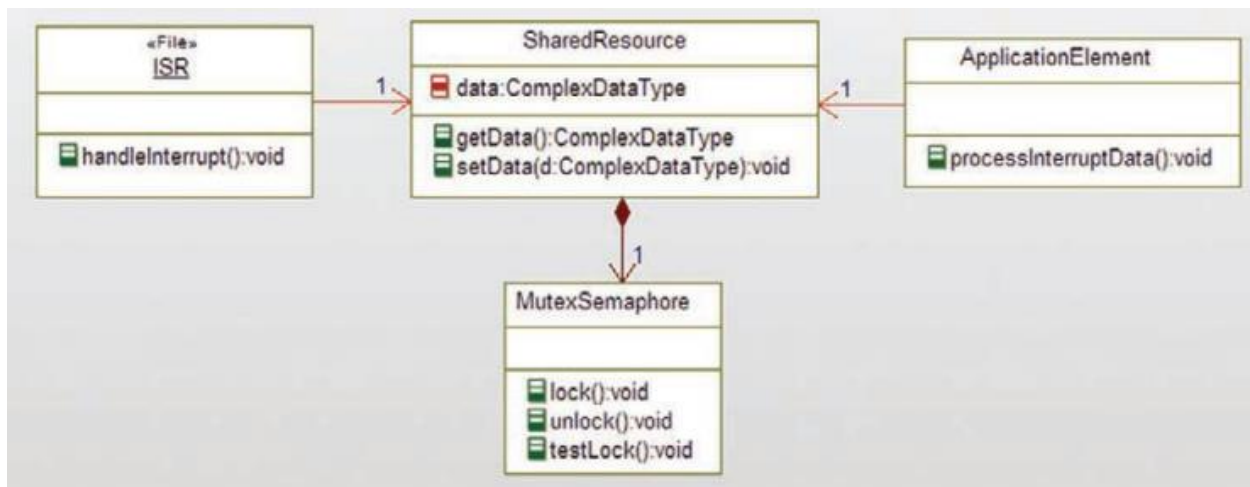
Consequences

- This pattern allows for **highly responsive processing of events** of interest.
- Normally, **interrupts are disabled while an interrupt service routine is executing**; this means that interrupt service routines must execute **very quickly** to ensure that other interrupts are not missed.
- Problems arise with this pattern when the ISR processing takes too long, when an implementation **mistake leaves interrupts disabled**, or **race conditions** or **deadlocks** occur on shared resources.

Potential Race Condition



Potential Deadlock Condition



Polling Pattern

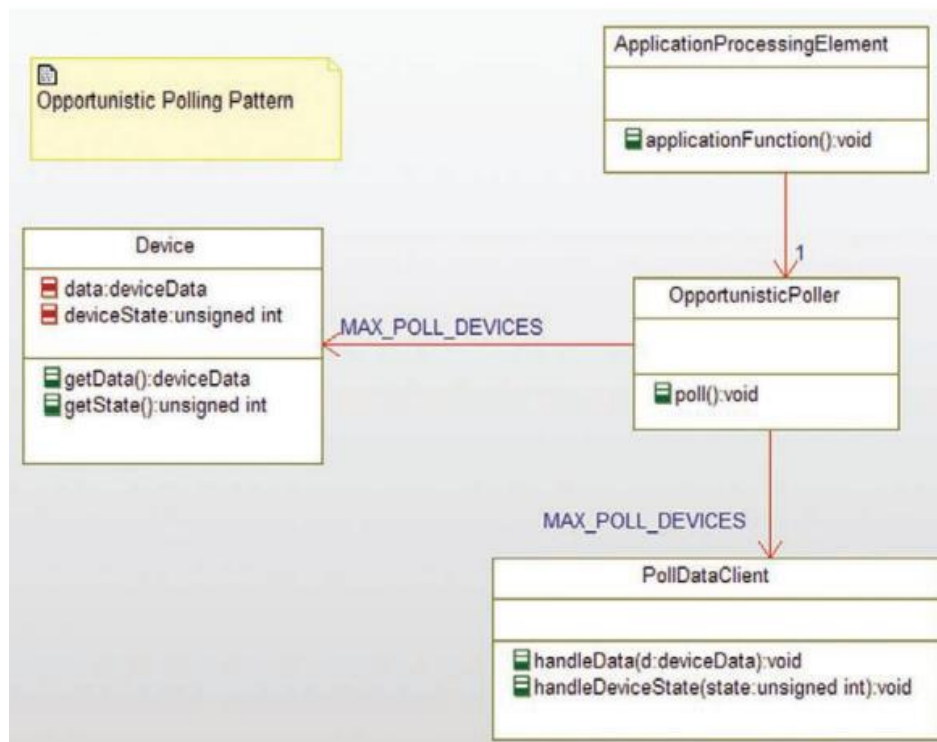
Abstract

- A common pattern for getting sensor data or signals from hardware is to **check periodically**, a process known as **polling**.
- Polling is **useful** when the data or signals are **not so urgent that they can wait** until the next polling period to be received or when the hardware isn't capable of generating interrupts when data or signals become available.
- Polling can be **periodic** or **opportunistic (waiting for the right moment)**.
- periodic polling uses a timer to indicate when the hardware should be sampled
- opportunistic polling is done when it is convenient for the system

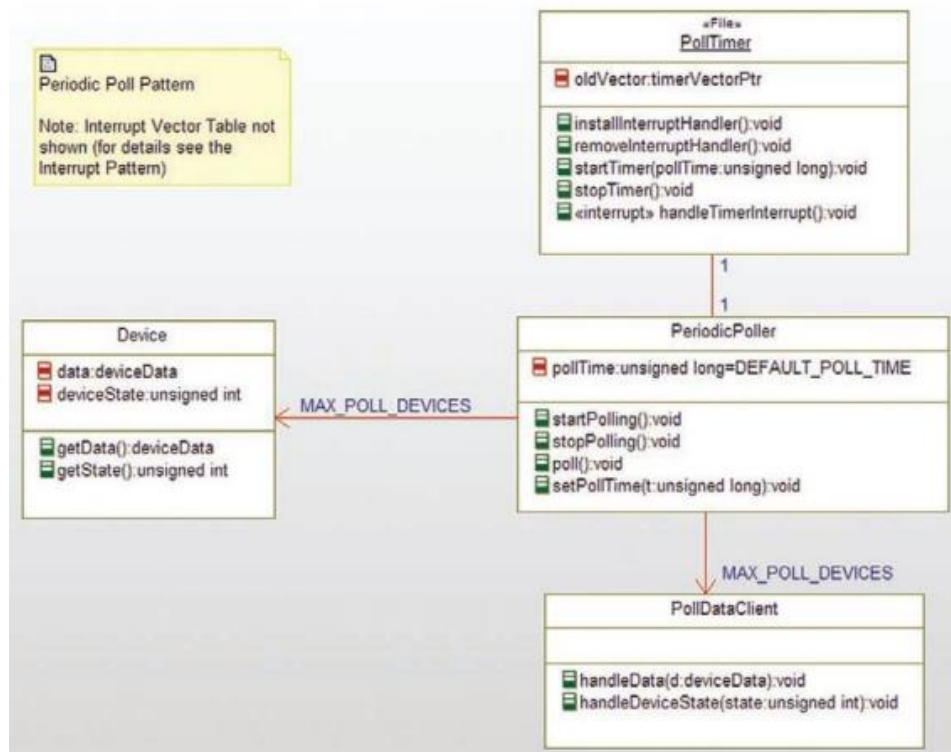
Problem

- The Polling Pattern addresses the concern of getting new sensor data or hardware signals into the system as it runs when the data or events are **not highly urgent** and the **time** between data sampling **can be guaranteed to be fast enough**.

Opportunistic Polling Pattern Structure



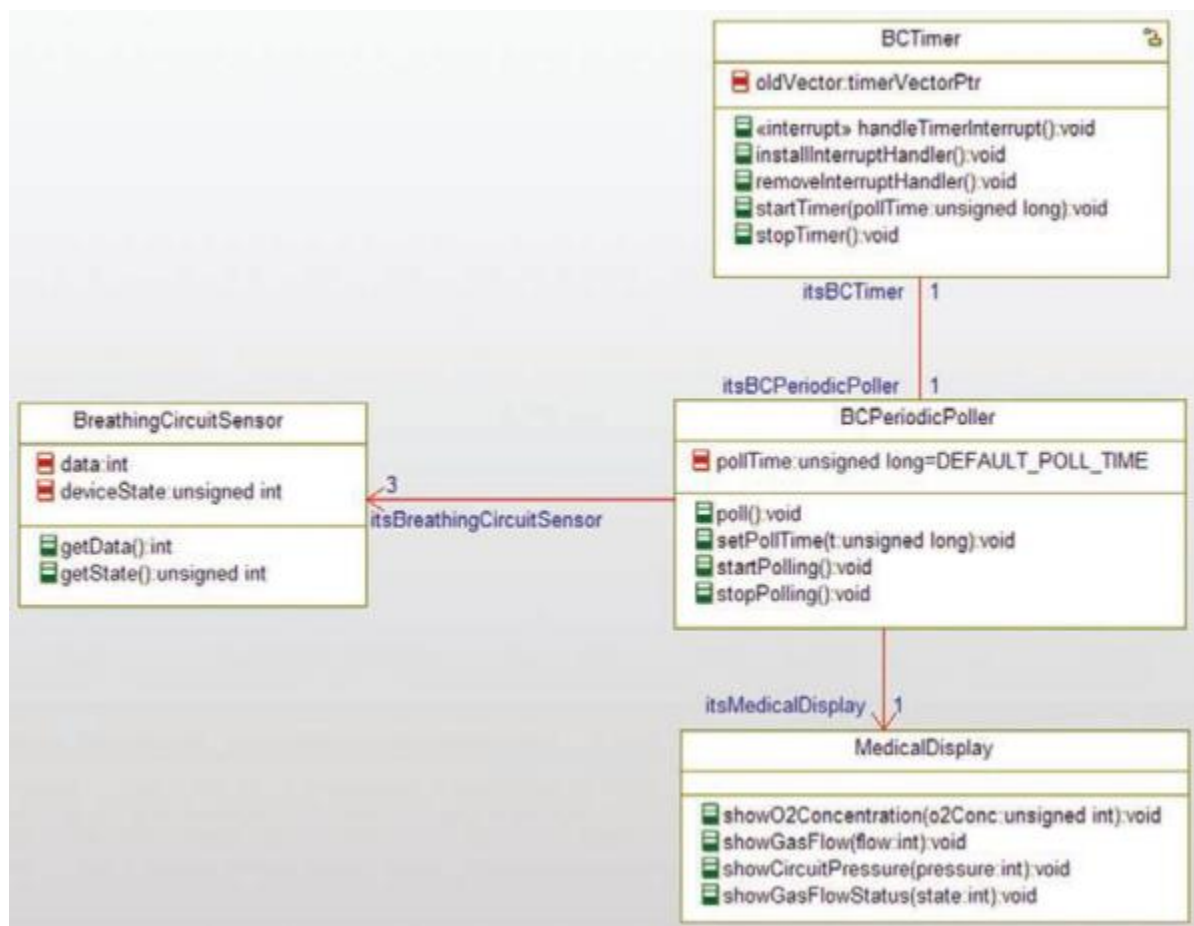
Periodic Polling Pattern Structure



Consequences

- Polling is **simpler** than the setup and use of the ISRs, although periodic polling is usually implemented with an ISR tied to a poll timer.
- Polling can **check many different devices at the same time** for status changes but is usually less timely than interrupts.
- If data **arrive faster** than the poll time, then data will be **lost**.
- This is not a problem in many applications but is **fatal in others**.

Example



Lecture 9

● Cyclic Executive Pattern

Abstract

- used for very small embedded applications where it allows multiple tasks to be run pseudo-concurrently without the overhead of a complex scheduler or RTOS.

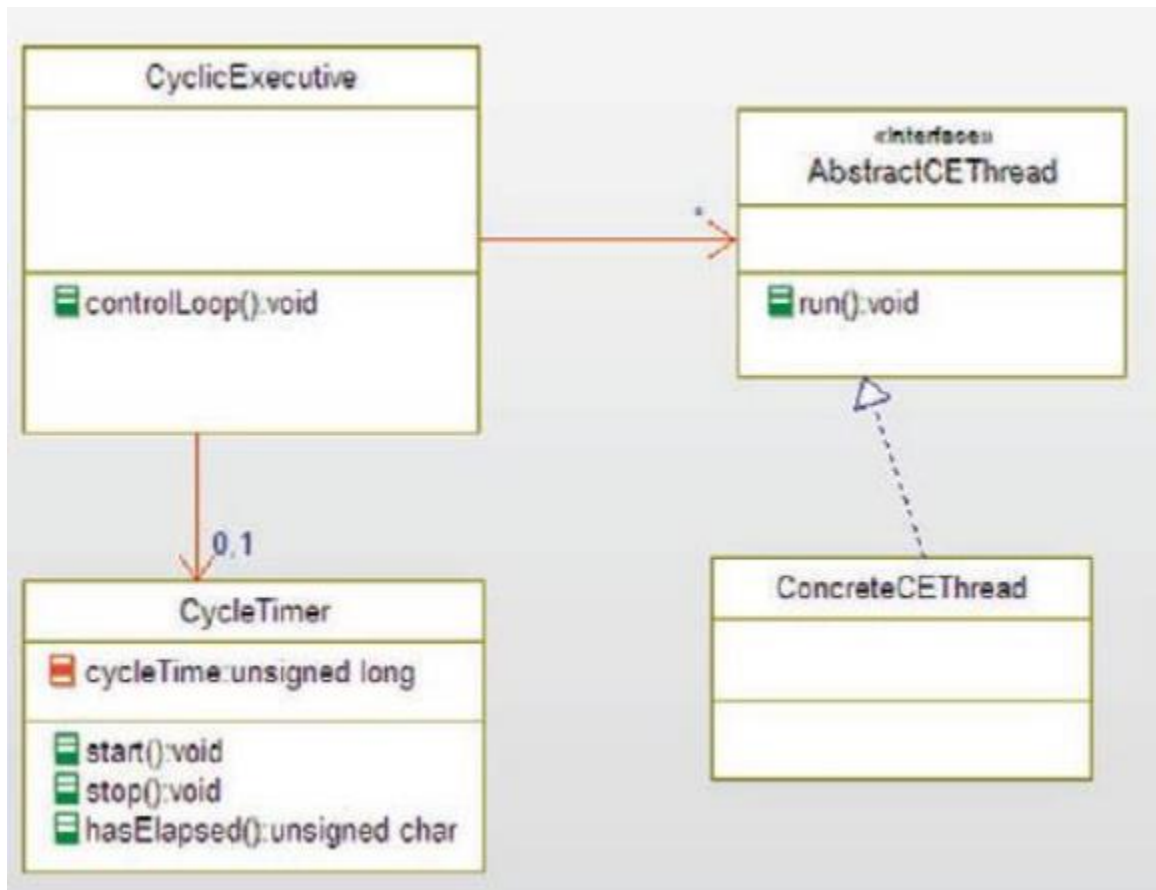
Problem

- Many embedded systems are tiny applications that have extremely tight memory and time constraints and cannot possibly host even a scaled-down microkernel RTOS.
- The Cyclic Executive Pattern provides a low-resource means to accomplish Schedulability for a set of tasks.
 - D_j is the deadline for task j
 - C_j is the worst-case execution time for task j
 - K is the cyclic executive loop overhead, including the overhead of the task invocation and return

for all tasks i , $D_i \geq \sum_{j=1}^n C_j + K$

Equation 4-1: Schedulability for Cyclic Executive with n tasks

Structure

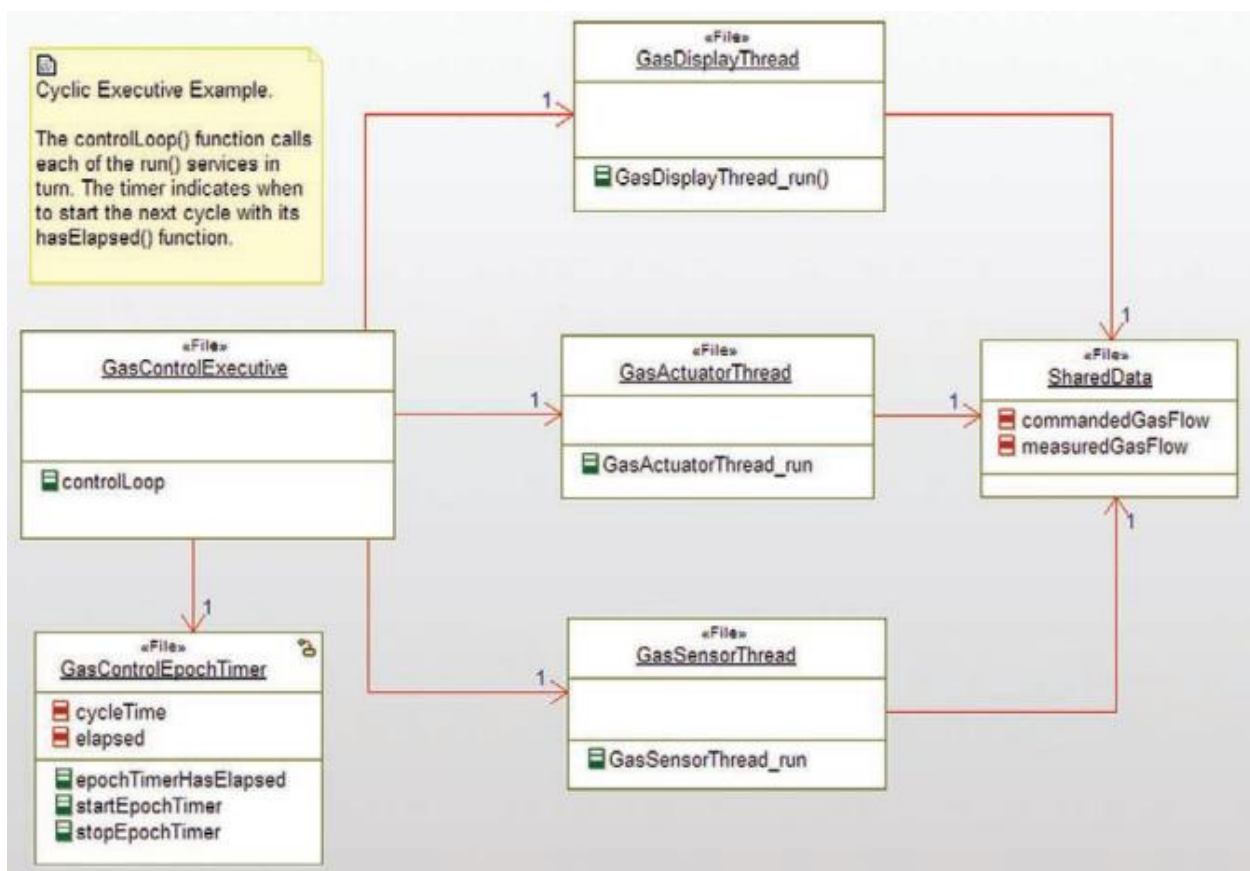


- The **AbstractCEThread** provides a standard interface for the threads by declaring a `run()` function.
- It is this function that the Cyclic Executive will **invoke** to execute the task.
- When this function completes, the CyclicExecutive runs the next task in the list.
- The **CyclicExecutive** contains the infinite loop that runs the tasks each in turn.
- The **CyclicExecutive** contains **global stack and static data that are needed** either by the tasks.
- The **time-triggered Cyclic Executive** will set up and use the **CycleTimer** to initiate each **epoch (execution of the task list)**.
- because of the **lack of preemption, unbounded blocking is not a concern**.

Consequences

- The pattern is **less responsive to high-urgency events**, and this makes the pattern inappropriate for applications that have high responsiveness requirements.
- A downside of using the pattern is that the **interaction of threads is made more challenging** than other patterns.
- If an output of one task is needed by another, **the data must be retained in global memory or a shared resource until that task runs**

Example



● Critical Region Pattern

Abstract

- The critical region pattern used for **task coordination** around.
- This happens by **disabling task switching during a region** where it is important that the current task executes without being preempted.
- This can be because it is **accessing a resource that cannot be simultaneously accessed safely**, or because it is crucial that the task executes for a portion of time without interruption.

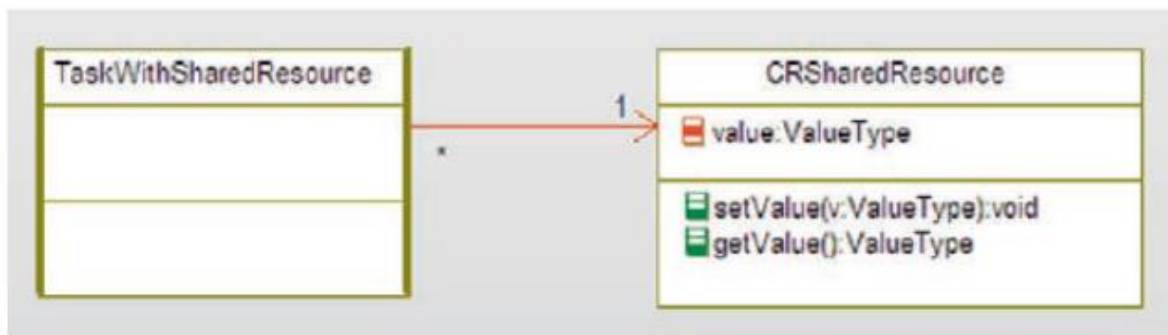
Problem

- There are two primary circumstances under which a task should not be interrupted or preempted.
- First, it may be accessing **a resource that may not be safely accessed by multiple clients simultaneously**.
- **Secondly, the task may be performing actions that must be completed within a short period of time** or that must be performed in a specific order.
- This pattern allows the active task to execute without any potential interference from other tasks.

Structure



Figure 4-12: Critical Region Pattern



Consequences

- This pattern enforces the critical region policies well and can be used to **turn off the scheduled task switches or, disable interrupts all together**.
- Care must be taken to **reenable task switching once the element is out of its critical region**.
- Additionally, the use of this pattern can **affect the timing of other tasks**. Critical regions are **usually short in duration** for this reason.
- There is **no problem of unbounded priority inversion** because all **task switching is disabled** during the critical region.
- Care must be taken when **nesting calls** to functions with their own critical regions.
- If such a function is called within the critical region of another, **the nest function call will reenale task switching and end the critical region** of the caller function.