



**MENOUFIA UNIVERSITY**  
**FACULTY OF COMPUTERS AND INFORMATION**

**Fourth Year (Second Semester)**  
**CS Dept., (CS 436 )**

---

---

# **Natural Language Processing**

## **NLP**

---

---

**Lecture Two**

**Dr. Hamdy M. Mousa**

# Regular Expression

- Regular Expression (**RE**)
  - (first developed by Kleene (1956)
  - is **a formula** in a special language that is used for specifying simple classes of strings.
  - is **a pattern** that matches some sequence in a text.
  - It is a mixture of:
    - characters or strings of text
    - special characters
    - groups or ranges
- Since common text-processing programs agree on most of the syntax of regular expressions,
  - all UNIX, **Microsoft Word**, and WordPerfect regular expressions.

# Regular Expression

- Regular expression search requires a pattern that we want to search for.
  - a corpus of texts to search through.
- A regular expression search function will search through the corpus returning all texts that contain the pattern.
  - In an information retrieval (IR) system such as a web search engine, the texts might be entire documents or web pages.
  - In a word-processor, the texts might be individual words, or lines of a document.
  - A search can be designed to return all matches to a regular expression or only the first match.

# Basic Regular Expression Patterns

- The simplest kind of regular expression is a sequence of simple characters.
  - Example:
    - To search for **woodchuck**, we type /woodchuck/.
  - So the regular expression /Buttercup/ matches any string containing the substring Buttercup,
    - Example: I'm called little Buttercup

## **NOTE:**

we will put slashes around each regular expression to make it clear

# Regular Expressions

RE	Example Patterns Matched
/woodchucks/	“interesting links to <u>woodchucks</u> and lemurs”
/a/	“M <u>a</u> ry Ann stopped by Mona’s”
/Claire_says,/	“Dagmar, my gift please,” <u>Claire</u> says,”
/song/	“all our pretty <u>songs</u> ”
/!/	“You’ve left the burglar behind again <u>!</u> ” said Nori

- Regular expressions are **case sensitive**.

# The brackets [ ]

The use of the brackets [ ] to specify a **disjunction** of characters.

RE	Match	Example Patterns
/ [wW] oodchuck /	Woodchuck or woodchuck	“ <u>W</u> oodchuck”
/ [abc] /	‘a’, ‘b’, <i>or</i> ‘c’	“In uomini, in soldat <u>i</u> ”
/ [1234567890] /	any digit	“plenty of <u>7</u> to 5”

The use of the brackets [ ] **plus the dash** - to specify any one character in a range.

RE	Match	Example Patterns Matched
/ [A-Z] /	an uppercase letter	“we should call it ‘ <u>D</u> renched Blossoms”
/ [a-z] /	a lowercase letter	“ <u>m</u> y beans were impatient to be hoed!”
/ [0-9] /	a single digit	“Chapter <u>1</u> : Down the Rabbit Hole”

# caret (^)

- The brackets can also be used to specify what a single character cannot be, by use of the caret (^).
- If the caret ^ is the **first symbol** after the open brackets [ , the resulting pattern is **negated**.

**Example:** the pattern / [ ^a] / matches any single character (including special characters) except **a**.

RE	Match (single characters)	Example Patterns Matched
[ ^A-Z]	not an uppercase letter	“Oy <u>fn</u> pripetchik”
[ ^Ss]	neither ‘S’ nor ‘s’	“ <u>I</u> have no exquisite reason for’t”
[ ^\.]	not a period	“ <u>o</u> ur resident Djinn”
[ e^]	either ‘e’ or ‘^’	“look up <u>^</u> now”
a^b	the pattern ‘a^b’	“look up <u>a^b</u> now”

# Question-mark `/?/`

- Optional elements
- How do we specify both `woodchuck` and `woodchucks`?
- The question-mark `/?/`, means “the preceding character or nothing”.

RE	Match	Example Patterns Matched
<code>woodchucks?</code>	woodchuck or woodchucks	<u>“woodchuck”</u>
<code>colou?r</code>	color or colour	<u>“colour”</u>



# Kleene \*

- The Kleene \* (pronounced “cleany star”) means 'zero or more occurrences of the immediately previous character or regular expression'.
  - So /a\*/ means 'any string of zero or more a's'.
    - This will match a or aaaaaa
  - Write RE for: An integer (a string of digits)?

# Kleene +

- This is the Kleene +, which means 'one or more of the previous character'.
- Thus the expression / [ 0- 9 ] + / is the normal way to specify 'a sequence of digits'.

# The period (/./)

- One very important **special character** is the period (/./, a wildcard expression that matches **any single character** (except a carriage return).

RE	Match	Example Patterns
/beg.n/	any character between 'beg' and 'n'	<u>begin</u> , <u>beg'n</u> , <u>begun</u>

# Anchors (caret ^)

- **Anchors** are special characters that anchor regular expressions to particular places in a string.
  - The most common anchors are the caret ^ and the dollar-sign \$.
  - The caret ^ matches the start of a line.
    - The pattern / ^ The/ matches the word **The** only at the start of a line.

# Anchors (dollar-sign \$)

- The most common anchors are **the caret ^ and the dollar-sign \$**.
  - The caret ^ matches the **start** of a line
  - The dollar sign \$ matches the **end of a line** (a space at the end of a line)

**Ex.:** / ^ **The dog**\. \$ / matches a line that contains only the phrase ***The dog***.

- use the backslash here since we want the **.** to mean 'period' .

# Anchors (\b & \B)

- There are also two other anchors:
  - \b matches a word boundary,
  - \B matches a non-boundary.
  - Thus \bthe\b matches the word *the* but not the word *other*.
- Guess , \b99/ will match ???

# Special characters

- Special characters for **start** and **end**:
- `/^man/` => any sequence which begins with “man”: *man, manned, manning...*
- `/man$/` => any sequence ending with “man”: *human, policeman ...*
- `/^man$/` => any sequence consisting of “man” only

# Disjunction, Grouping, and Precedence

- In such a case, we might want to search for either the string *cat* or the string *dog*.
  - we need a new operator, **the disjunction** operator, also called the pipe symbol **|**.
  - The pattern `/cat | dog/` matches either the string *cat* or the string *dog*.



# Disjunction, Grouping, and Precedence

- For example, How can I specify both guppy and guppies?
  - Cannot say **/ guppy | ies/**, because that would match only the strings *guppy* and *ies*.
  - This is because sequences like *guppy* take precedence over the disjunction operator **|**.
- So the pattern **/gupp (y|ies) /** would specify that we meant the disjunction only to apply to the suffixes *y* and *ies*.

# Quantifiers

- **{m}**
  - Specifies that exactly ***m* copies** of the previous RE should be matched.
    - **Ex.:** `a{6}` will match exactly six 'a' characters, but not five.
- **{m, n}**
  - Causes the resulting RE to match from *m* to *n* repetitions of the preceding RE.
    - For example, `a{3,5}` will match from 3 to 5 'a' characters.
  - Omitting *m* specifies a lower bound of zero,
  - Omitting *n* specifies an infinite upper bound.
    - **Ex.:** `a{4,} b` will match `aaaab` or a thousand 'a' characters followed by a `b`, but not `aaab`.

# Quantifiers

- `/ba*/`
  - matches *b*, *ba*, *baa*, *baaa*
  - `/*` means “zero or more of the preceding character or group”
- `/(ba ){1,3}/`
  - matches *ba*, *ba ba* or *ba ba ba*
  - `{n, m}` means “between n and m”
- `/(ba ){2}/`
  - matches *ba ba*
  - `{n}` means “exactly n”

# Operator Precedence

- |                        |                |
|------------------------|----------------|
| 1. Parentheses         | ( )            |
| 2. Counters            | * + ? { }      |
| 3. Sequence of Anchors | the ^my end \$ |
| 4. Disjunction         |                |

- **Example:**

- /moo+/
- /try|ies/
- /and|or/

# Exercise

- Write a regular expression to find all instances of the determiner “*the*”:

*The recent attempt by the police to retain their current rates of pay has not gathered much favor with the southern factions.*

# Exercise

- **/the/**

*The recent attempt by the police to retain their current rates of pay has not gathered much favor with the southern factions.*

- **/[Tt]he/**

*The recent attempt by the police to retain their current rates of pay has not gathered much favor with the southern factions.*

- **/b[Tt]he\b/**

*The recent attempt by the police to retain their current rates of pay has not gathered much favor with the southern factions.*

# Advanced Operators

RE	Expansion	Match	Example Patterns
\d	[0-9]	any digit	Party_of_5
\D	[^0-9]	any non-digit	Blue_moon
\w	[a-zA-Z0-9_]	any alphanumeric or space	Daiyu
\W	[^\w]	a non-alphanumeric	!!!!
\s	[\r\t\n\f]	whitespace (space, tab)	
\S	[^\s]	Non-whitespace	in_Concord

- Finally, certain special characters are referred to by special notation **based on the backslash (/).**
  - the newline character /n and the tab character /t.
- To refer to characters that are special themselves, (like., \*, [, and/), precede them with a backslash, (i.e. /\./, /\\*/ , /\[/, and /\//).

# Regular Expression

## Regular expression operators for counting

RE	Match
*	zero or more occurrences of the previous char or expression
+	one or more occurrences of the previous char or expression
?	exactly zero or one occurrence of the previous char or expression
{ n }	n occurrences of the previous char or expression
{ n , m }	from n to m occurrences of the previous char or expression
{ n , }	at least n occurrences of the previous char or expression

## Some characters that need to be backslashed

RE	Match	Example Patterns Matched
\ *	an asterisk “*”	“K_A*P*L*A*N”
\ .	a period “.”	“Dr. Livingston, I presume”
\ ?	a question mark	“Would you light my candle?”
\ n	a newline	
\ t	a tab	



# Exercise

- Write RE to represent fractions of dollars.  
(\$199.99, ....)
- Write RE to represent:  
“any computer with more than 6 GHz and  
500 GB of disk space for less than \$1000”.

# Regular Expression Substitution

- An important use of regular expressions is in substitutions.
- We'd like a way to refer back to the integer, we've found so that we can easily add the brackets.
  - To do this, we put parentheses ( and) around the first pattern,
  - and use the number operator \1 in the second pattern to refer back. Here's how it looks:
- `s/( [0-9] +)/<\1>/`
- **Ex.:** changing : 35 to <35>

# Regular Expression Substitution

- The parenthesis and number operators can also be used to specify that a certain string or expression must occur twice in the text.

i.e. : suppose we are looking for the pattern  
'the Xer they were, the Xer they will be

## Example :

/the (.\*) er they (.\*) , the \1er they \2/

The bigger they were, the bigger they were

but not **The bigger they were, the bigger they will be.**

□ These numbered memories are called registers

# Non-capturing Group

- Occasionally we might want to use **parentheses** for grouping, but **don't** want to capture the resulting pattern in a register.
  - use a **non-capturing group**, which is specified by putting the commands group **?:** after the open parenthesis.

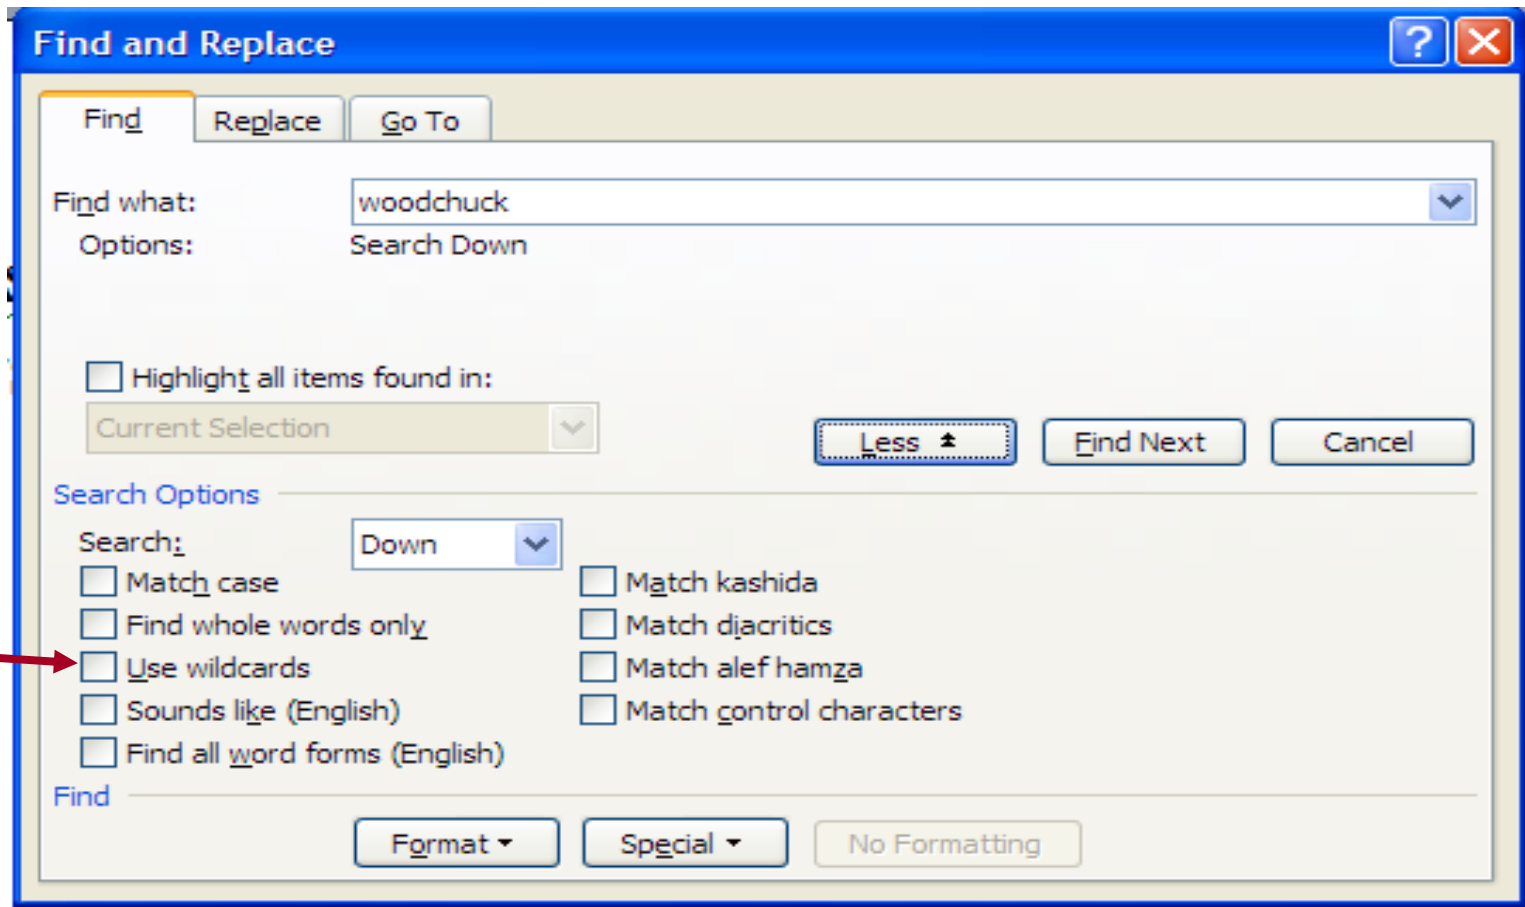
**(?: pattern )**

- Ex.:**

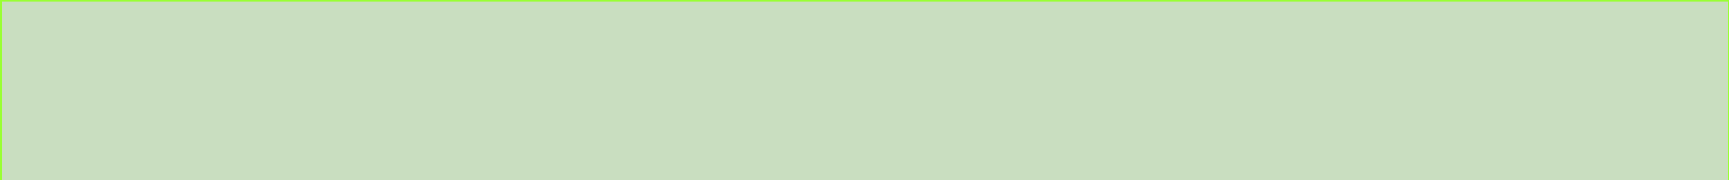
`/(?:some|a few) (people|cats) like some \1/`

# Report

- Try regular expressions in MS WORD in both Arabic & English



- Write RE can match "gray" or "grey"
  - gray|grey
  - gr(a|e)y
- $a|b^*$  and  $(a|b)^*$  are equivalent
  - $a|b^*$  denotes  $\{\epsilon, "a", "b", "bb", "bbb", \dots\}$
  - $(a|b)^*$  denotes the set of all strings with no symbols other than "a" and "b", including the empty string:  $\{\epsilon, "a", "b", "aa", "ab", "ba", "bb", "aaa", \dots\}$

- 
- Normalizing text means converting it to a more convenient, standard form
  - Another part of text normalization is lemmatization, the task of determining that two words have the same root, despite their surface differences
  - Stemming refers to a simpler version of lemmatization in which we mainly just strip suffixes from the end of the word.